

notebook_final

April 23, 2021

1 Final Project Submission

- Student name: Jonathan Lee
- Student pace: full time
- Scheduled project review date/time: April 27, 2pm
- Instructor name: James Irving

1.1 TABLE OF CONTENTS

Click to jump to matching Markdown Header.

- **Introduction**
- **OBTAIN**
- **SCRUB**
- **EXPLORE**
- **MODEL**
- **iNTERPRET**
- **Conclusions/Recommendations** _____

2 INTRODUCTION

This analysis focuses on creating a multiple regression model based on housing data from King County, Washington. We will work through an exploratory data analysis to clean the data that we have to prepare it for modeling, as well as working through an iterative approach to refining our model. The goal of this analysis is to create a model which explains how different attributes affect the value of a housing property in King County, and to extract specific variables which we can use to recommend to a homeowner in King County how to increase the value of his/her home.

3 OBTAIN

The data that we will use in this analysis has been provided as a .csv file. We will inspect the data types to determine how to approach the cleansing process.

```
[1]: # Import packages to be used in notebook.  
import pandas as pd  
import numpy as np  
import seaborn as sns
```

```

import matplotlib.pyplot as plt
from matplotlib.ticker import FuncFormatter
from matplotlib.gridspec import GridSpec

import statsmodels.api as sm
import statsmodels.stats.api as sms
import statsmodels.formula.api as smf

from scipy import stats

from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import StandardScaler

import warnings
warnings.simplefilter(action="ignore", category=FutureWarning)

%matplotlib inline

```

```

[2]: # Load housing data
df = pd.read_csv('data/kc_house_data.csv')
display(df.head(5), df.info())

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     21597 non-null  int64
1   date                   21597 non-null  object
2   price                  21597 non-null  float64
3   bedrooms               21597 non-null  int64
4   bathrooms              21597 non-null  float64
5   sqft_living             21597 non-null  int64
6   sqft_lot                21597 non-null  int64
7   floors                  21597 non-null  float64
8   waterfront              19221 non-null  float64
9   view                    21534 non-null  float64
10  condition               21597 non-null  int64
11  grade                   21597 non-null  int64
12  sqft_above              21597 non-null  int64
13  sqft_basement           21597 non-null  object
14  yr_built                21597 non-null  int64
15  yr_renovated            17755 non-null  float64
16  zipcode                 21597 non-null  int64
17  lat                     21597 non-null  float64
18  long                    21597 non-null  float64
19  sqft_living15           21597 non-null  int64

```

```

20 sqft_lot15      21597 non-null int64
dtypes: float64(8), int64(11), object(2)
memory usage: 3.5+ MB

```

	id	date	price	bedrooms	bathrooms	sqft_living \
0	7129300520	10/13/2014	221900.0	3	1.00	1180
1	6414100192	12/9/2014	538000.0	3	2.25	2570
2	5631500400	2/25/2015	180000.0	2	1.00	770
3	2487200875	12/9/2014	604000.0	4	3.00	1960
4	1954400510	2/18/2015	510000.0	3	2.00	1680

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement \
0	5650	1.0	NaN	0.0	...	7	1180	0.0
1	7242	2.0	0.0	0.0	...	7	2170	400.0
2	10000	1.0	0.0	0.0	...	6	770	0.0
3	5000	1.0	0.0	0.0	...	7	1050	910.0
4	8080	1.0	0.0	0.0	...	8	1680	0.0

	yr_built	yr_renovated	zipcode	lat	long	sqft_living15	sqft_lot15
0	1955	0.0	98178	47.5112	-122.257	1340	5650
1	1951	1991.0	98125	47.7210	-122.319	1690	7639
2	1933	NaN	98028	47.7379	-122.233	2720	8062
3	1965	0.0	98136	47.5208	-122.393	1360	5000
4	1987	0.0	98074	47.6168	-122.045	1800	7503

```
[5 rows x 21 columns]
```

```
None
```

4 SCRUB

The data looks clean for the most part, but there are null values in the columns labeled 'waterfront', 'view' and 'yr_renovated' which will be addressed in this section. We also need to make sure to address the two columns that have been stored as object data types labeled 'date' and 'sqft_basement' in addition to checking for duplicated entries.

4.1 Checking for duplicates

```
[3]: df[df['id'].duplicated(keep=False)]
```

```

# Duplicates in id have different dates, and can be considered as resold
↳ properties.

```

```

[3]:
      id      date      price  bedrooms  bathrooms  sqft_living \
93  6021501535  7/25/2014  430000.0         3         1.50         1580
94  6021501535 12/23/2014  700000.0         3         1.50         1580
313 4139480200  6/18/2014 1380000.0         4         3.25         4290
314 4139480200 12/9/2014 1400000.0         4         3.25         4290

```

324	7520000520	9/5/2014	232000.0	2	1.00	1240
...
20654	8564860270	3/30/2015	502000.0	4	2.50	2680
20763	6300000226	6/26/2014	240000.0	4	1.00	1200
20764	6300000226	5/4/2015	380000.0	4	1.00	1200
21564	7853420110	10/3/2014	594866.0	3	3.00	2780
21565	7853420110	5/4/2015	625000.0	3	3.00	2780

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	\
93	5000	1.0	0.0	0.0	...	8	1290	
94	5000	1.0	0.0	0.0	...	8	1290	
313	12103	1.0	0.0	3.0	...	11	2690	
314	12103	1.0	0.0	3.0	...	11	2690	
324	12092	1.0	NaN	0.0	...	6	960	
...	
20654	5539	2.0	NaN	0.0	...	8	2680	
20763	2171	1.5	0.0	0.0	...	7	1200	
20764	2171	1.5	0.0	0.0	...	7	1200	
21564	6000	2.0	0.0	0.0	...	9	2780	
21565	6000	2.0	0.0	0.0	...	9	2780	

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	\
93	290.0	1939	0.0	98117	47.6870	-122.386	
94	290.0	1939	0.0	98117	47.6870	-122.386	
313	1600.0	1997	0.0	98006	47.5503	-122.102	
314	1600.0	1997	0.0	98006	47.5503	-122.102	
324	280.0	1922	1984.0	98146	47.4957	-122.352	
...	
20654	0.0	2013	0.0	98045	47.4759	-121.734	
20763	0.0	1933	0.0	98133	47.7076	-122.342	
20764	0.0	1933	0.0	98133	47.7076	-122.342	
21564	0.0	2013	0.0	98065	47.5184	-121.886	
21565	0.0	2013	NaN	98065	47.5184	-121.886	

	sqft_living15	sqft_lot15
93	1570	4500
94	1570	4500
313	3860	11244
314	3860	11244
324	1820	7460
...
20654	2680	5992
20763	1130	1598
20764	1130	1598
21564	2850	6000
21565	2850	6000

[353 rows x 21 columns]

```
[4]: df[df.duplicated(keep=False)]  
  
# There are no duplicated entries
```

```
[4]: Empty DataFrame  
Columns: [id, date, price, bedrooms, bathrooms, sqft_living, sqft_lot, floors,  
waterfront, view, condition, grade, sqft_above, sqft_basement, yr_built,  
yr_renovated, zipcode, lat, long, sqft_living15, sqft_lot15]  
Index: []  
  
[0 rows x 21 columns]
```

```
[5]: # Drop id and date columns since they are not controllable attributes  
# that would affect the value of the property  
df.drop(['id', 'date'], axis=1, inplace=True)
```

4.2 Checking null value counts

```
[6]: # Check number of NaN cells in dataframe  
df.isna().sum()
```

```
[6]: price                0  
bedrooms                0  
bathrooms              0  
sqft_living            0  
sqft_lot               0  
floors                 0  
waterfront            2376  
view                   63  
condition              0  
grade                  0  
sqft_above             0  
sqft_basement          0  
yr_built               0  
yr_renovated          3842  
zipcode                0  
lat                    0  
long                   0  
sqft_living15          0  
sqft_lot15             0  
dtype: int64
```

```
[7]: # Check waterfront column's value counts  
df['waterfront'].value_counts(dropna=False)
```

```
[7]: 0.0    19075
      NaN    2376
      1.0     146
      Name: waterfront, dtype: int64
```

We will assume that homes with a missing value for ‘waterfront’ are not located on a waterfront

```
[8]: # Replace Nan cells with 0.0
      df['waterfront'].fillna(0.0, inplace=True)
```

```
[9]: # Confirm that fillna method worked properly
      df['waterfront'].value_counts(dropna=False)
```

```
[9]: 0.0    21451
      1.0     146
      Name: waterfront, dtype: int64
```

```
[10]: #Check yr_renovated value counts
      df['yr_renovated'].value_counts(dropna=False)
```

```
[10]: 0.0    17011
      NaN    3842
      2014.0     73
      2003.0     31
      2013.0     31
      ...
      1944.0      1
      1948.0      1
      1976.0      1
      1934.0      1
      1953.0      1
      Name: yr_renovated, Length: 71, dtype: int64
```

Similar to the ‘waterfront’, we will assume that homes with a missing value for ‘yr_renovated’ have not undergone renovation.

```
[11]: # Replace NaN cells with 0.0
      df['yr_renovated'].fillna(0.0, inplace=True)
```

```
[12]: # Confirm that fillna method worked properly
      df['yr_renovated'].value_counts(dropna=False)
```

```
[12]: 0.0    20853
      2014.0     73
      2003.0     31
      2013.0     31
      2007.0     30
      ...
```

```

1946.0      1
1959.0      1
1971.0      1
1951.0      1
1954.0      1
Name: yr_renovated, Length: 70, dtype: int64

```

Due to the ambiguous definition of the ‘view’ column, we will drop it to avoid including any variables in our regression model that we cannot explain.

```

[13]: # Remove view due to ambiguous definition
df.drop('view', axis=1, inplace=True)

```

```

[14]: #Confirm that all NaN cells have been addressed
df.isna().sum()

```

```

[14]: price      0
      bedrooms    0
      bathrooms    0
      sqft_living    0
      sqft_lot      0
      floors      0
      waterfront    0
      condition    0
      grade        0
      sqft_above    0
      sqft_basement 0
      yr_built     0
      yr_renovated  0
      zipcode      0
      lat         0
      long        0
      sqft_living15 0
      sqft_lot15   0
      dtype: int64

```

4.3 Converting Data Types

Great, no more null values to address. Now we need to check why ‘sqft_basement’ is being stored as an object data type. We will go ahead and remove the missing entries since the count is not large and convert the data type to float or int.

```

[15]: # Check for non-numerical entries in sqft_basement
df['sqft_basement'].value_counts().sort_index()

```

```

[15]: 0.0      12826
      10.0       2
      100.0     42

```

```

1000.0      148
1008.0       1
...
960.0       65
970.0       44
980.0       57
990.0       52
?          454
Name: sqft_basement, Length: 304, dtype: int64

```

```

[16]: # Remove entries where sqft_basement is '?'
df = df[df['sqft_basement'] != '?']

```

```

[17]: # Convert sqft_basement from object to float
df['sqft_basement'] = df['sqft_basement'].astype(float)

```

4.4 Feature Engineering

Because those properties that have not been renovated contain a value of 0.0 under their 'yr_renovated' column, this will skew the rest of the data where the other entries that have been renovated will contain a year number. We will engineer a binary feature that indicates whether or not the property has undergone any renovation in order to avoid this skew issue.

```

[18]: # Define function to create column with value 1 if renovated, 0 if not
      ↪renovated.
def renov_bool(row):
    if row['yr_renovated'] > 0:
        val = 1
    else:
        val = 0
    return val

```

```

[19]: # Apply function to create renovated column
df['renovated'] = df.apply(renov_bool, axis=1)
df.drop('yr_renovated', axis=1, inplace=True)

```

```

[20]: # Verify that we have successfully removed null values and fixed data types
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 21143 entries, 0 to 21596
Data columns (total 18 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   price           21143 non-null  float64
 1   bedrooms        21143 non-null  int64
 2   bathrooms       21143 non-null  float64

```



```

3  sqft_living      21143 non-null  int64
4  sqft_lot         21143 non-null  int64
5  floors           21143 non-null  float64
6  waterfront       21143 non-null  float64
7  condition        21143 non-null  int64
8  grade            21143 non-null  int64
9  sqft_above       21143 non-null  int64
10 sqft_basement    21143 non-null  float64
11 yr_built         21143 non-null  int64
12 zipcode          21143 non-null  int64
13 lat              21143 non-null  float64
14 long             21143 non-null  float64
15 sqft_living15    21143 non-null  int64
16 sqft_lot15       21143 non-null  int64
17 renovated        21143 non-null  int64
dtypes: float64(7), int64(11)
memory usage: 3.1 MB

```

4.5 Checking for Correlation and Multicollinearity

We will move on to check for how correlated each column is with our target variable 'price' as well as check for multicollinearity

```

[21]: # Create correlation matrix from dataframe
price_corr = df.corr()
price_corr.round(2)

```

```

[21]:
price      price  bedrooms  bathrooms  sqft_living  sqft_lot  floors  \
price      1.00    0.31    0.53    0.70    0.09    0.26
bedrooms   0.31    1.00    0.51    0.58    0.03    0.18
bathrooms  0.53    0.51    1.00    0.76    0.09    0.50
sqft_living 0.70    0.58    0.76    1.00    0.17    0.35
sqft_lot    0.09    0.03    0.09    0.17    1.00   -0.01
floors      0.26    0.18    0.50    0.35   -0.01    1.00
waterfront  0.27    0.00    0.06    0.11    0.02    0.02
condition   0.04    0.03   -0.13   -0.06   -0.01   -0.26
grade       0.67    0.36    0.67    0.76    0.11    0.46
sqft_above  0.61    0.48    0.69    0.88    0.18    0.52
sqft_basement 0.33    0.30    0.28    0.43    0.02   -0.25
yr_built    0.05    0.16    0.51    0.32    0.05    0.49
zipcode     -0.05   -0.15   -0.20   -0.20   -0.13   -0.06
lat          0.31   -0.01    0.02    0.05   -0.09    0.05
long         0.02    0.13    0.22    0.24    0.23    0.13
sqft_living15 0.59    0.39    0.57    0.76    0.14    0.28
sqft_lot15    0.08    0.03    0.09    0.18    0.72   -0.01
renovated     0.12    0.02    0.05    0.05    0.00    0.00

waterfront  condition  grade  sqft_above  sqft_basement  \

```

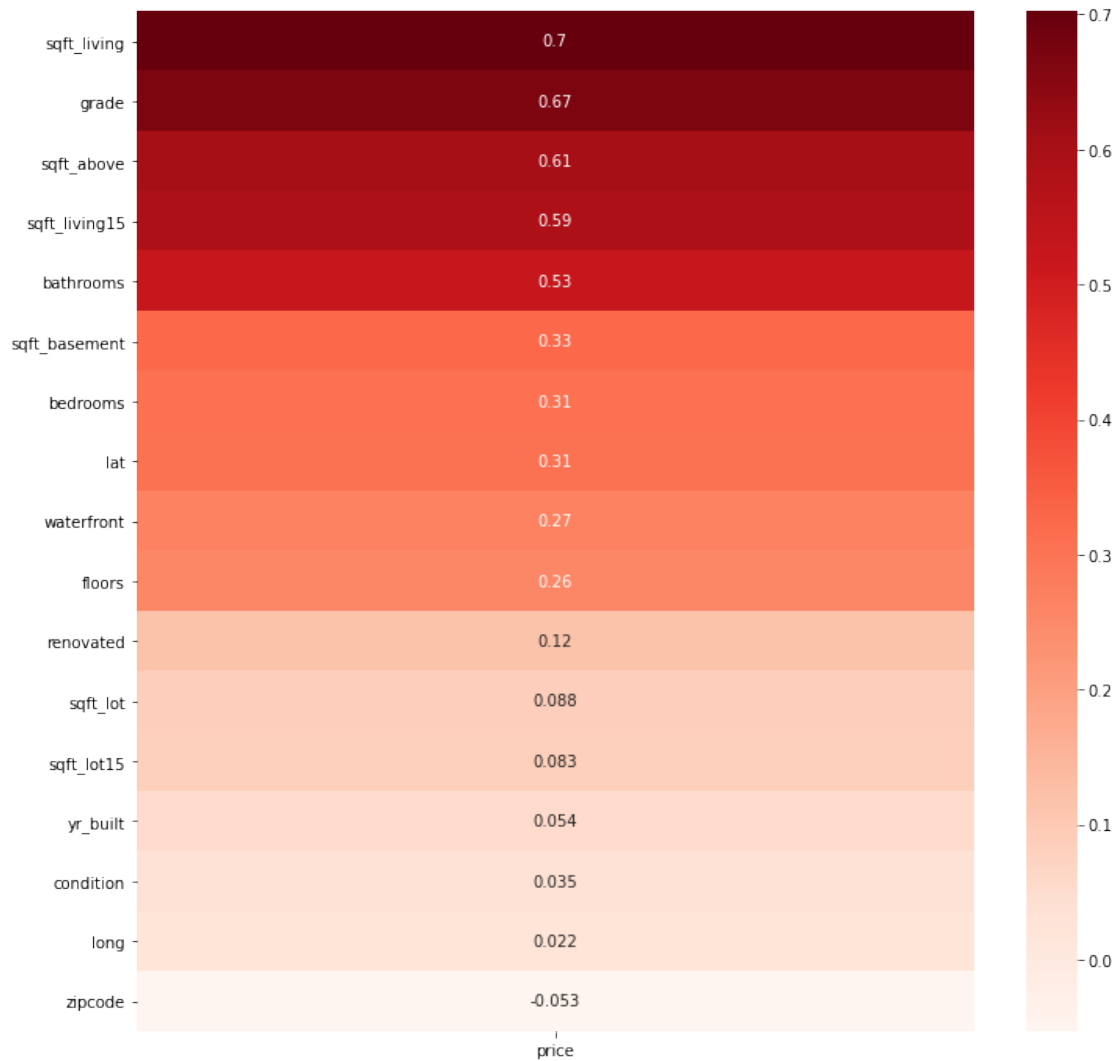
price	0.27	0.04	0.67	0.61	0.33
bedrooms	0.00	0.03	0.36	0.48	0.30
bathrooms	0.06	-0.13	0.67	0.69	0.28
sqft_living	0.11	-0.06	0.76	0.88	0.43
sqft_lot	0.02	-0.01	0.11	0.18	0.02
floors	0.02	-0.26	0.46	0.52	-0.25
waterfront	1.00	0.02	0.08	0.07	0.08
condition	0.02	1.00	-0.15	-0.16	0.17
grade	0.08	-0.15	1.00	0.76	0.17
sqft_above	0.07	-0.16	0.76	1.00	-0.05
sqft_basement	0.08	0.17	0.17	-0.05	1.00
yr_built	-0.02	-0.36	0.45	0.43	-0.13
zipcode	0.03	0.00	-0.19	-0.26	0.08
lat	-0.01	-0.02	0.11	-0.00	0.11
long	-0.04	-0.11	0.20	0.34	-0.15
sqft_living15	0.09	-0.09	0.71	0.73	0.20
sqft_lot15	0.03	-0.00	0.12	0.20	0.02
renovated	0.07	-0.06	0.02	0.02	0.07

	yr_built	zipcode	lat	long	sqft_living15	sqft_lot15	\
price	0.05	-0.05	0.31	0.02	0.59	0.08	
bedrooms	0.16	-0.15	-0.01	0.13	0.39	0.03	
bathrooms	0.51	-0.20	0.02	0.22	0.57	0.09	
sqft_living	0.32	-0.20	0.05	0.24	0.76	0.18	
sqft_lot	0.05	-0.13	-0.09	0.23	0.14	0.72	
floors	0.49	-0.06	0.05	0.13	0.28	-0.01	
waterfront	-0.02	0.03	-0.01	-0.04	0.09	0.03	
condition	-0.36	0.00	-0.02	-0.11	-0.09	-0.00	
grade	0.45	-0.19	0.11	0.20	0.71	0.12	
sqft_above	0.43	-0.26	-0.00	0.34	0.73	0.20	
sqft_basement	-0.13	0.08	0.11	-0.15	0.20	0.02	
yr_built	1.00	-0.35	-0.15	0.41	0.33	0.07	
zipcode	-0.35	1.00	0.27	-0.56	-0.28	-0.15	
lat	-0.15	0.27	1.00	-0.14	0.05	-0.08	
long	0.41	-0.56	-0.14	1.00	0.33	0.26	
sqft_living15	0.33	-0.28	0.05	0.33	1.00	0.18	
sqft_lot15	0.07	-0.15	-0.08	0.26	0.18	1.00	
renovated	-0.20	0.06	0.03	-0.06	0.00	0.00	

	renovated
price	0.12
bedrooms	0.02
bathrooms	0.05
sqft_living	0.05
sqft_lot	0.00
floors	0.00
waterfront	0.07

condition	-0.06
grade	0.02
sqft_above	0.02
sqft_basement	0.07
yr_built	-0.20
zipcode	0.06
lat	0.03
long	-0.06
sqft_living15	0.00
sqft_lot15	0.00
renovated	1.00

```
[22]: # Correlation heatmap customization guide was utilized to create the following
      ↪ visualizations:
      # https://medium.com/@chrisshaw982/
      ↪ seaborn-correlation-heatmaps-customized-10246f4f7f4b
      fig, ax = plt.subplots(figsize=(12,12))
      sns.heatmap(price_corr[['price']].drop('price').sort_values(by='price',
      ↪ ascending=False), annot=True,
                  ax=ax, cmap='Reds');
```




```
[23]: # Drop price to only show correlation between independent variables
corr = df.drop('price', axis=1).corr().round(2)
```

```
[24]: # Create mask for upper triangle of matrix
mask = np.zeros_like(corr)

mask[np.triu_indices_from(mask)] = True
```

```
[25]: #Create heatmap correlation matrix
fig, ax = plt.subplots(figsize=(12,12))
sns.heatmap(corr, annot=True, ax=ax, cmap='coolwarm', vmin=-1, vmax=1,\
            mask=mask);
```




From the correlation heatmap, we can see that other than 'sqft_living', we do not have any variables that are high enough to remove prior to running our baseline model. We will go ahead and remove 'sqft_living' to address the issue of multicollinearity in our dataset.

```
[26]: # Remove sqft_living to get address multicollinearity
df.drop('sqft_living', axis=1, inplace=True)

corr = df.drop('price', axis=1).corr().round(2)

mask = np.zeros_like(corr)
mask[np.triu_indices_from(mask)] = True
```

```
[27]: # Check heatmap correlation matrix after removing column
fig, ax = plt.subplots(figsize=(12,12))
sns.heatmap(corr, annot=True, ax=ax, cmap='coolwarm', vmin=-1, vmax=1,
            mask=mask);
```



5 EXPLORE

In this section, we will explore the distributions as well as addressing the issue of outliers in each column. We will also be checking to see how much of a linear relationship each variable has with our target variable 'price'.

```
[28]: # Set theme and style for plots.
sns.set_theme('talk')
sns.set_style('darkgrid')
```

```
[29]: df.describe()
```

```
[29]:
```

	price	bedrooms	bathrooms	sqft_lot	floors	\
count	2.114300e+04	21143.000000	21143.000000	2.114300e+04	21143.000000	
mean	5.405107e+05	3.372558	2.116079	1.508714e+04	1.493591	
std	3.680751e+05	0.924917	0.768531	4.120920e+04	0.539249	
min	7.800000e+04	1.000000	0.500000	5.200000e+02	1.000000	
25%	3.220000e+05	3.000000	1.750000	5.043000e+03	1.000000	
50%	4.500000e+05	3.000000	2.250000	7.620000e+03	1.500000	
75%	6.450000e+05	4.000000	2.500000	1.069550e+04	2.000000	
max	7.700000e+06	33.000000	8.000000	1.651359e+06	3.500000	

	waterfront	condition	grade	sqft_above	sqft_basement	\
count	21143.000000	21143.000000	21143.000000	21143.000000	21143.000000	
mean	0.006716	3.409923	7.658279	1789.069006	291.851724	
std	0.081679	0.650498	1.174253	828.409769	442.498337	
min	0.000000	1.000000	3.000000	370.000000	0.000000	
25%	0.000000	3.000000	7.000000	1200.000000	0.000000	
50%	0.000000	3.000000	7.000000	1560.000000	0.000000	
75%	0.000000	4.000000	8.000000	2210.000000	560.000000	
max	1.000000	5.000000	13.000000	9410.000000	4820.000000	

	yr_built	zipcode	lat	long	sqft_living15	\
count	21143.000000	21143.000000	21143.000000	21143.000000	21143.000000	
mean	1971.023223	98077.868893	47.560274	-122.213876	1987.27139	
std	29.321938	53.535756	0.138591	0.140597	685.67034	
min	1900.000000	98001.000000	47.155900	-122.519000	399.00000	
25%	1952.000000	98033.000000	47.471250	-122.328000	1490.00000	
50%	1975.000000	98065.000000	47.572000	-122.230000	1840.00000	
75%	1997.000000	98117.000000	47.678200	-122.125000	2360.00000	
max	2015.000000	98199.000000	47.777600	-121.315000	6210.00000	

	sqft_lot15	renovated
count	21143.000000	21143.000000
mean	12738.941967	0.034196
std	27169.273663	0.181736
min	651.000000	0.000000
25%	5100.000000	0.000000
50%	7626.000000	0.000000
75%	10087.000000	0.000000
max	871200.000000	1.000000

5.1 Checking for Normality, Outliers, and Linearity

There appear to be some outliers, as in the case of bedrooms where the max number is 33. Although this might be an error in data collection, we will leave the outliers be for now to see how they affect the skew of our data and how our baseline model turns out with what has been provided.

We will proceed to visualize how our data is distributed as well as the linearity of each variable against the price variable.

```
[30]: # Create function to plot histogram and boxplot to indicate normality and
      ↪ outliers
      # and scatterplot to show linearity with the target variable
      def plot_distribution_linearity(df, col=None, verbose=False, boxplot=True):
          """This function was written by James Irving during study group.
          Original function has been modified to include regression plot to
          illustrate linear relationship with 'price' column.

          Plots a histogram + KDE and a boxplot of the column.
          Also prints statistics for skew, kurtosis, and normaltest.

          Args:
              df_ (DataFrame): DataFrame containing column to plot
              col (str): Name of the column to plot.
              verbose (bool, optional): If true show figure and print stats. Defaults
          ↪ to True.
              boxplot (bool, optional): If true, return subplots with boxplot.
          ↪ Defaults to True.

          Returns:
              fig : Matplotlib Figure
              ax : Matplotlib Axis
          """

          # df = df_.copy()

          if col is None:
              data = df.copy()
              name = data.name
          else:
              data = df[col].copy()
              name = col

          ## Calc mean and mean skew and curtosis
          median = data.median().round(2)
          mean = data.mean().round(2)
          skew_val = round(stats.skew(data, bias=False), 2)
          kurt_val = round(stats.kurtosis(data, bias=False), 2)
```



```

## Plot distribution
fig = plt.figure(figsize=(11, 6))
gs = GridSpec(nrows=2, ncols=2)

ax0 = fig.add_subplot(gs[0, 0])
ax1 = fig.add_subplot(gs[1, 0])
ax2 = fig.add_subplot(gs[:, 1])

sns.histplot(data,alpha=0.5,stat='density',ax=ax0)
sns.kdeplot(data,color='green',label='KDE',ax=ax0)
ax0.set(ylabel='Density',title=name.title())
ax0.set_title(F"Distribution of {name.title()}")
ax0.axvline(median,label=f'median={median:,.}',color='black')
ax0.axvline(mean,label=f'mean={mean:,.}',color='black',ls=':')
ax0.legend()

## Plot Boxplot
sns.boxplot(data,x=col,ax=ax1)
ax1.set_title(F"Box Plot of {name.title()}")

# Plot Scatterplot to illustrate linearity
sns.regplot(data=df, x=col, y='price', line_kws={"color": "red"}, ax=ax2)
ax2.set_title(F"Scatter Plot of {name.title()}")

## Tweak Layout & Display
fig.tight_layout()

## Delete boxplot if unwanted
if boxplot == False:
    fig.delaxes(ax[1])

if verbose:
    plt.show()

    print('[i] Distribution Stats:')
    print(f"\tSkew = {skew_val}")
    print(f"\tKurtosis = {kurt_val}")
    print(f"\tN = {len(data):,}")

## Test for normality
result = stats.normaltest(data)
print('\n',result)
if result[1]<.05:
    print('\t- p<.05: The distribution is NOT normally distributed.')
elif result[1] >=.05:

```

```

        print('\t- p>=.05: The distribution IS normally distributed')

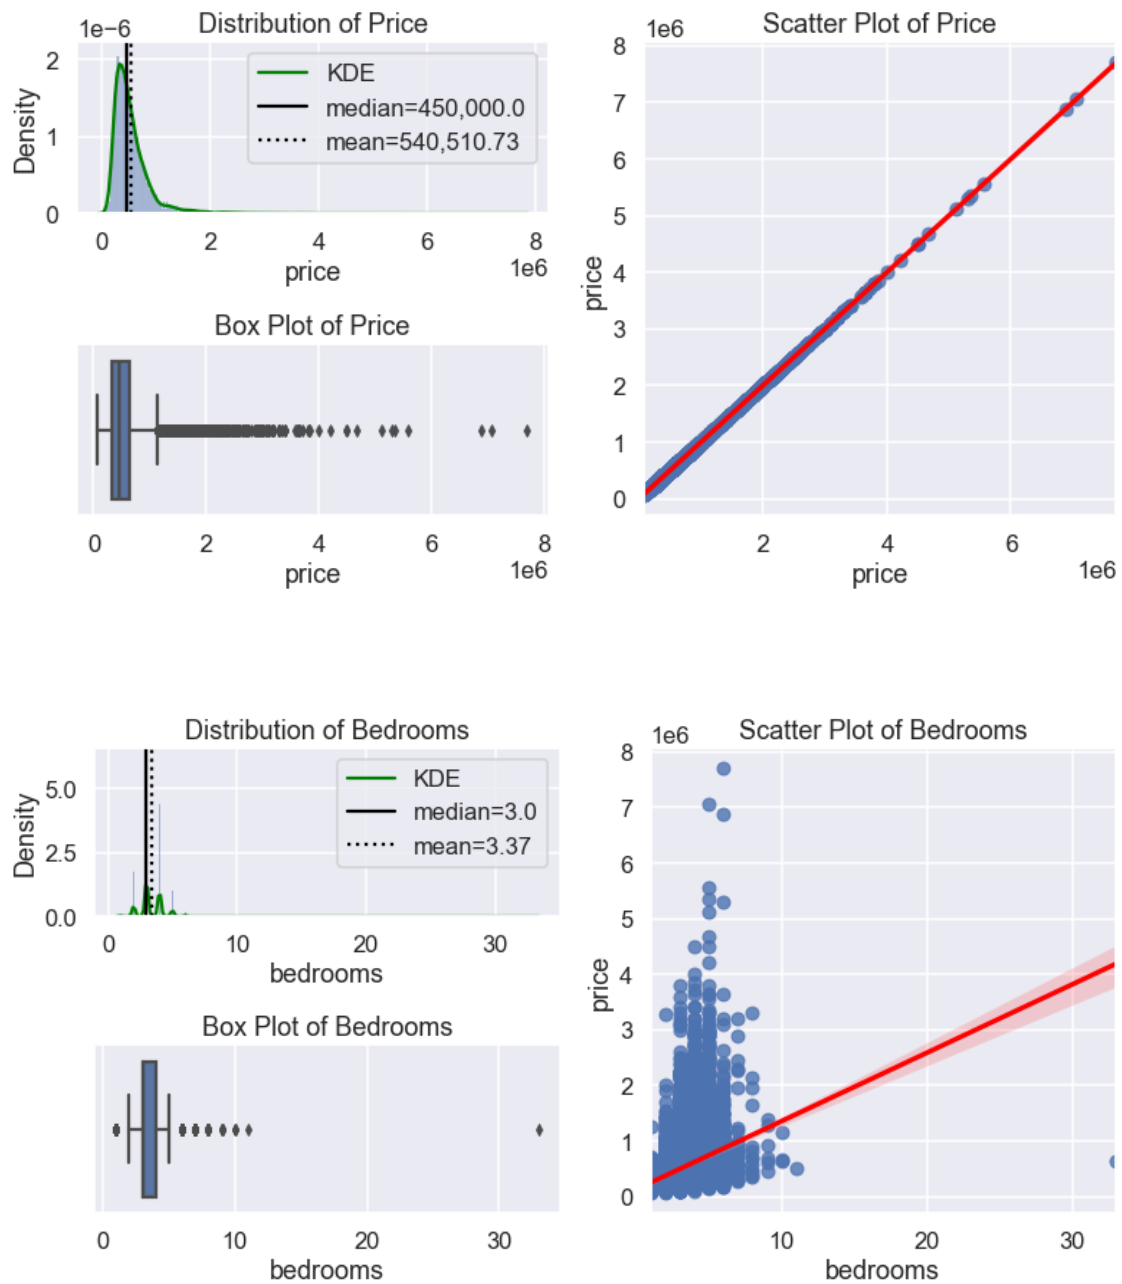
    return fig, ax




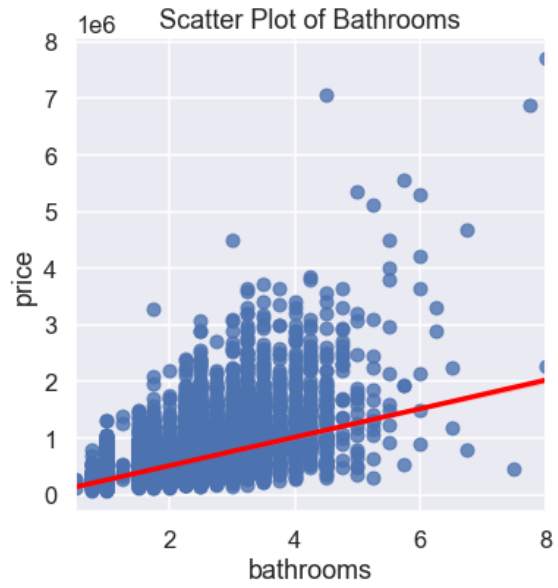

```

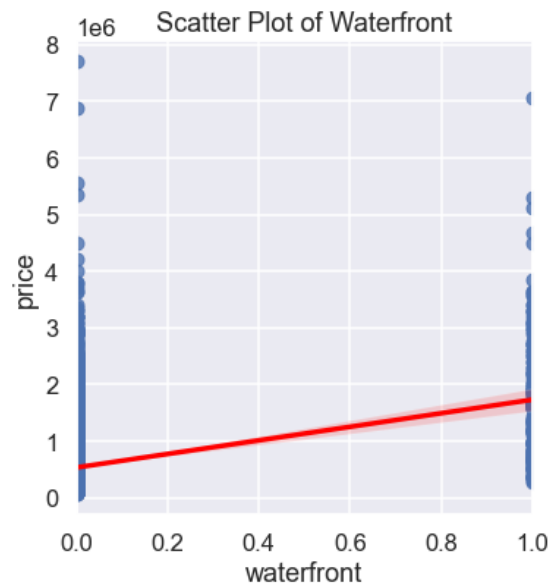


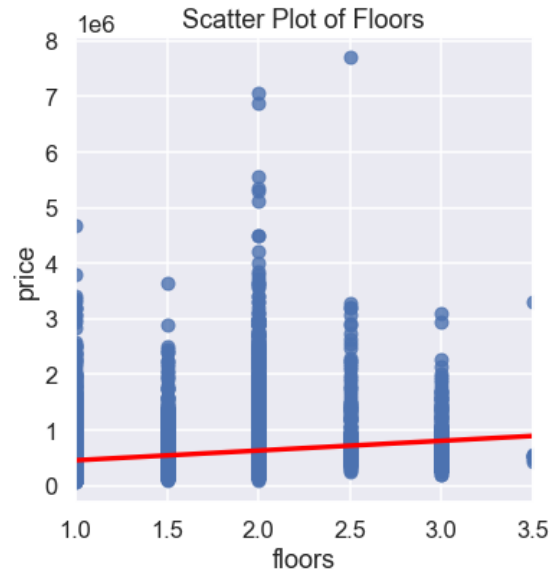


```




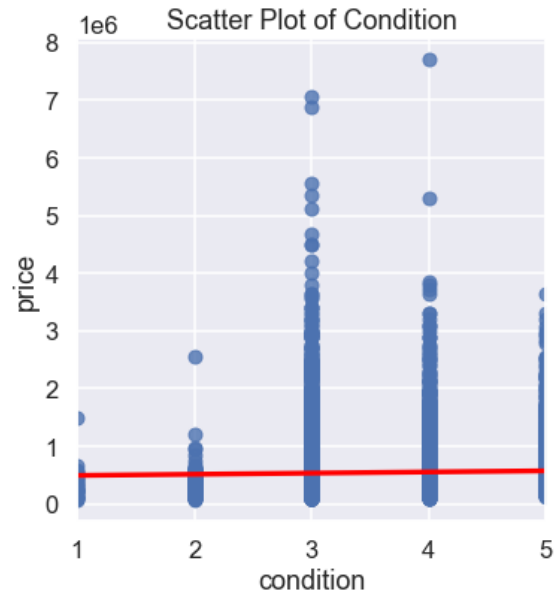


[31]: # Create plot for all columns
for col in df:
    plot_distribution_linearity(df=df, col=col);


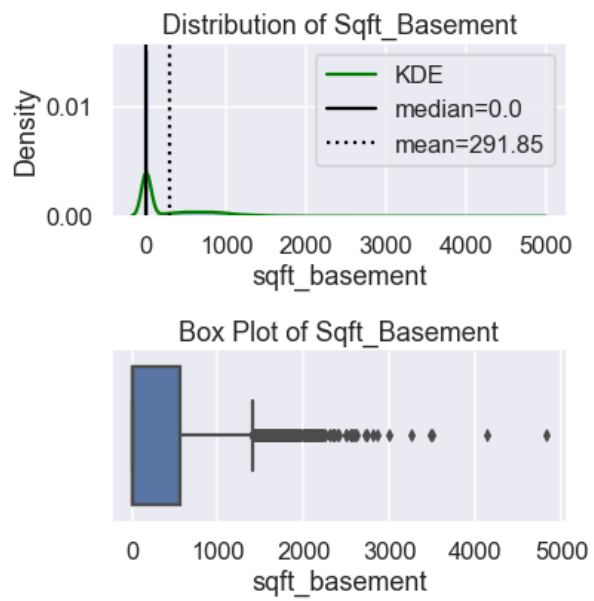

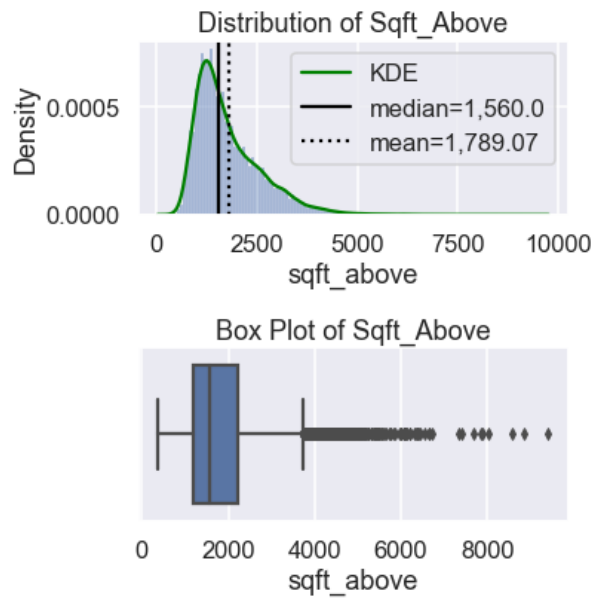
```









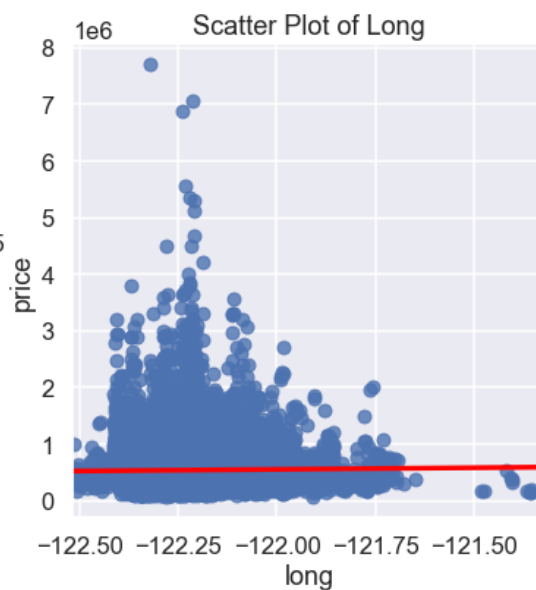
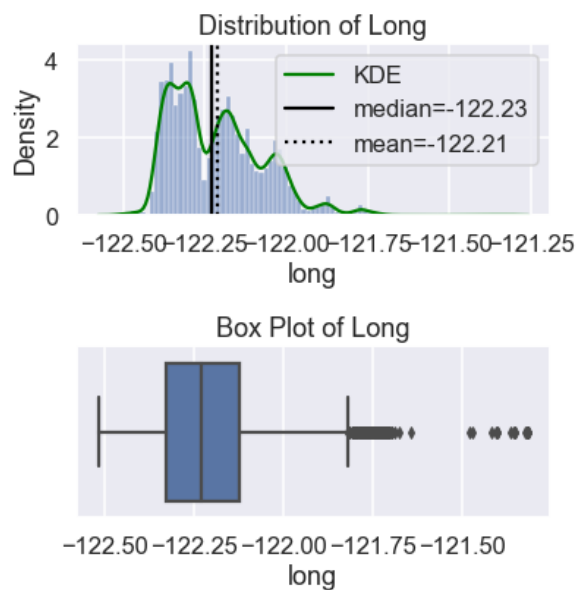












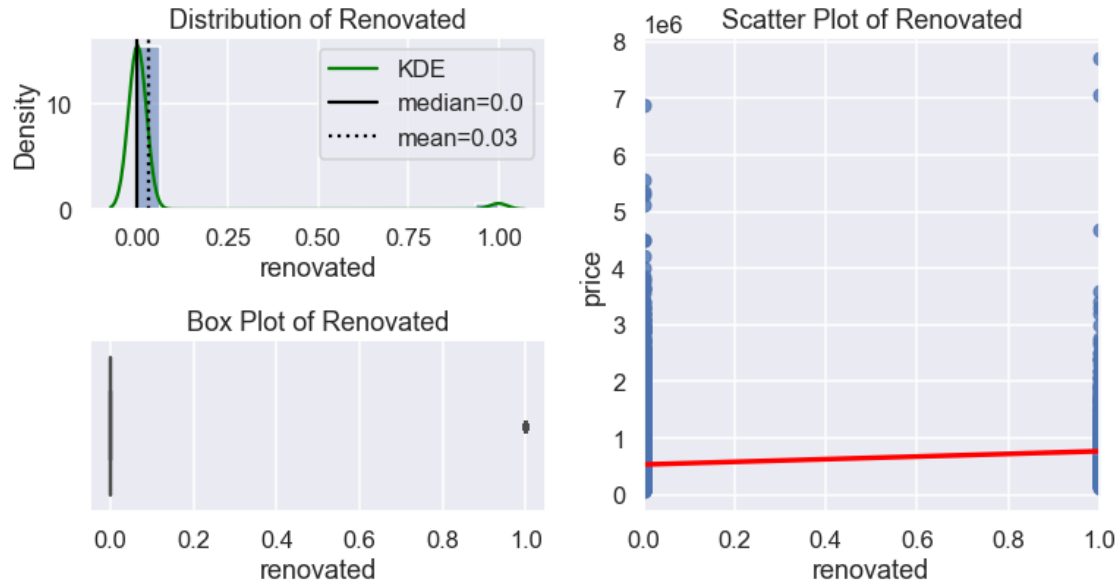












```
[32]: # Remove columns where there is weak linear relationship with price
df.drop(['condition', 'yr_built', 'renovated', 'sqft_lot15'], axis=1,
        inplace=True)
```

5.2 One Hot Encoding

We can see that there are some categorical variables in our dataset, but other than the 'zipcode' column, the other variables are ordinal.

We will proceed to use One Hot Encoding prior to running our multiple regression model including the zipcode data.

```
[33]: # One Hot Encode zipcodes column
encoder = OneHotEncoder(drop='first', sparse=False)
encoder.fit(df[['zipcode']])

ohe_vars = encoder.transform(df[['zipcode']])
ohe_vars

encoder.get_feature_names(['zipcode'])

df_ohe = pd.DataFrame(ohe_vars, columns=encoder.get_feature_names(['zipcode']),
                      index=df.index)
df_ohe
```

```
[33]:      zipcode_98002  zipcode_98003  zipcode_98004  zipcode_98005  \
0                0.0                0.0                0.0                0.0
1                0.0                0.0                0.0                0.0
```

2	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0
...
21592	0.0	0.0	0.0	0.0
21593	0.0	0.0	0.0	0.0
21594	0.0	0.0	0.0	0.0
21595	0.0	0.0	0.0	0.0
21596	0.0	0.0	0.0	0.0

	zipcode_98006	zipcode_98007	zipcode_98008	zipcode_98010	\
0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	
...	
21592	0.0	0.0	0.0	0.0	
21593	0.0	0.0	0.0	0.0	
21594	0.0	0.0	0.0	0.0	
21595	0.0	0.0	0.0	0.0	
21596	0.0	0.0	0.0	0.0	

	zipcode_98011	zipcode_98014	...	zipcode_98146	zipcode_98148	\
0	0.0	0.0	...	0.0	0.0	
1	0.0	0.0	...	0.0	0.0	
2	0.0	0.0	...	0.0	0.0	
3	0.0	0.0	...	0.0	0.0	
4	0.0	0.0	...	0.0	0.0	
...	
21592	0.0	0.0	...	0.0	0.0	
21593	0.0	0.0	...	1.0	0.0	
21594	0.0	0.0	...	0.0	0.0	
21595	0.0	0.0	...	0.0	0.0	
21596	0.0	0.0	...	0.0	0.0	

	zipcode_98155	zipcode_98166	zipcode_98168	zipcode_98177	\
0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	
...	
21592	0.0	0.0	0.0	0.0	
21593	0.0	0.0	0.0	0.0	
21594	0.0	0.0	0.0	0.0	
21595	0.0	0.0	0.0	0.0	

21596	0.0	0.0	0.0	0.0
	zipcode_98178	zipcode_98188	zipcode_98198	zipcode_98199
0	1.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0
...
21592	0.0	0.0	0.0	0.0
21593	0.0	0.0	0.0	0.0
21594	0.0	0.0	0.0	0.0
21595	0.0	0.0	0.0	0.0
21596	0.0	0.0	0.0	0.0

[21143 rows x 69 columns]

```
[34]: # Join One Hot Encoded dataframe with original dataframe and drop
# original zipcodes column
df_model = pd.concat([df.drop('zipcode',axis=1),df_ohe],axis=1)
df_model
```

```
[34]:
```

	price	bedrooms	bathrooms	sqft_lot	floors	waterfront	grade	\
0	221900.0	3	1.00	5650	1.0	0.0	7	
1	538000.0	3	2.25	7242	2.0	0.0	7	
2	180000.0	2	1.00	10000	1.0	0.0	6	
3	604000.0	4	3.00	5000	1.0	0.0	7	
4	510000.0	3	2.00	8080	1.0	0.0	8	
...		
21592	360000.0	3	2.50	1131	3.0	0.0	8	
21593	400000.0	4	2.50	5813	2.0	0.0	8	
21594	402101.0	2	0.75	1350	2.0	0.0	7	
21595	400000.0	3	2.50	2388	2.0	0.0	8	
21596	325000.0	2	0.75	1076	2.0	0.0	7	

	sqft_above	sqft_basement	lat	...	zipcode_98146	zipcode_98148	\
0	1180	0.0	47.5112	...	0.0	0.0	
1	2170	400.0	47.7210	...	0.0	0.0	
2	770	0.0	47.7379	...	0.0	0.0	
3	1050	910.0	47.5208	...	0.0	0.0	
4	1680	0.0	47.6168	...	0.0	0.0	
...	
21592	1530	0.0	47.6993	...	0.0	0.0	
21593	2310	0.0	47.5107	...	1.0	0.0	
21594	1020	0.0	47.5944	...	0.0	0.0	
21595	1600	0.0	47.5345	...	0.0	0.0	
21596	1020	0.0	47.5941	...	0.0	0.0	

	zipcode_98155	zipcode_98166	zipcode_98168	zipcode_98177	\
0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	
...	
21592	0.0	0.0	0.0	0.0	
21593	0.0	0.0	0.0	0.0	
21594	0.0	0.0	0.0	0.0	
21595	0.0	0.0	0.0	0.0	
21596	0.0	0.0	0.0	0.0	

	zipcode_98178	zipcode_98188	zipcode_98198	zipcode_98199
0	1.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0
...
21592	0.0	0.0	0.0	0.0
21593	0.0	0.0	0.0	0.0
21594	0.0	0.0	0.0	0.0
21595	0.0	0.0	0.0	0.0
21596	0.0	0.0	0.0	0.0

[21143 rows x 81 columns]

6 MODEL

Finally, we have prepared our data enough to be able to run an initial iteration of our multiple regression model! As we create each model, we will include a QQ plot to address the normality of residuals as well as plotting price vs residuals in order to check for homoscedasticity of residuals.

6.1 Creating a Baseline Model

```
[35]: # Create function to simultaneously run model and plot for normality
# and homoscedasticity of residuals.
def model_combined(df):
    ## Create a string representing the right side of the ~ in our formula
    features = ' + '.join(df.drop('price',axis=1).columns)

    ## Create the final formula and create the model
    f = "price~"+features
```

```

# Model regression
model = smf.ols(f, df).fit()
display(model.summary())

# Create QQ plot
fig, ax = plt.subplots(ncols=2,figsize=(14,6))
sm.graphics.qqplot(model.resid,dist=stats.norm,fit=True,line='45',\
                    ax=ax[0])
ax[0].set_title('QQ Plot')

# Create homoscedasticity plot
resids = model.resid
sns.scatterplot(x=model.predict(df.drop('price',axis=1), transform=True),\
                y=model.resid, ax=ax[1])
ax[1].axhline(0, color='r')
ax[1].set_title('Homoscedasticity of Residuals')
ax[1].set_xlabel('Predicted Price')
ax[1].set_ylabel('Residuals')

return model, fig, ax

```

```

[36]: # Run regression on cleaned dataframe
model_combined(df_model);

```

```

<class 'statsmodels.iolib.summary.Summary'>
"""

```

```

                                OLS Regression Results
=====
Dep. Variable:                  price    R-squared:                  0.793
Model:                            OLS    Adj. R-squared:              0.792
Method:                 Least Squares    F-statistic:                 1006.
Date:                Thu, 22 Apr 2021    Prob (F-statistic):           0.00
Time:                22:29:30    Log-Likelihood:             -2.8434e+05
No. Observations:          21143    AIC:                        5.689e+05
Df Residuals:              21062    BIC:                        5.695e+05
Df Model:                   80
Covariance Type:            nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-2.693e+07	6.51e+06	-4.140	0.000	-3.97e+07	-1.42e+07
bedrooms	-2.781e+04	1614.774	-17.224	0.000	-3.1e+04	-2.46e+04
bathrooms	1.29e+04	2621.916	4.920	0.000	7760.985	1.8e+04
sqft_lot	0.2365	0.031	7.642	0.000	0.176	0.297
floors	-6.476e+04	3127.358	-20.708	0.000	-7.09e+04	-5.86e+04
waterfront	8.813e+05	1.46e+04	60.424	0.000	8.53e+05	9.1e+05

grade	5.264e+04	1832.891	28.717	0.000	4.9e+04	5.62e+04
sqft_above	219.3557	3.115	70.429	0.000	213.251	225.460
sqft_basement	157.4838	3.686	42.720	0.000	150.258	164.709
lat	1.207e+05	6.69e+04	1.804	0.071	-1.04e+04	2.52e+05
long	-1.703e+05	4.83e+04	-3.523	0.000	-2.65e+05	-7.55e+04
sqft_living15	27.4453	2.985	9.193	0.000	21.594	33.297
zipcode_98002	5.715e+04	1.52e+04	3.760	0.000	2.74e+04	8.69e+04
zipcode_98003	-1.635e+04	1.37e+04	-1.195	0.232	-4.32e+04	1.05e+04
zipcode_98004	7.574e+05	2.47e+04	30.600	0.000	7.09e+05	8.06e+05
zipcode_98005	2.806e+05	2.64e+04	10.614	0.000	2.29e+05	3.32e+05
zipcode_98006	2.751e+05	2.16e+04	12.709	0.000	2.33e+05	3.18e+05
zipcode_98007	2.345e+05	2.73e+04	8.586	0.000	1.81e+05	2.88e+05
zipcode_98008	2.612e+05	2.6e+04	10.058	0.000	2.1e+05	3.12e+05
zipcode_98010	1.146e+05	2.33e+04	4.917	0.000	6.89e+04	1.6e+05
zipcode_98011	6.726e+04	3.38e+04	1.991	0.047	1034.942	1.33e+05
zipcode_98014	1.215e+05	3.71e+04	3.277	0.001	4.88e+04	1.94e+05
zipcode_98019	7.459e+04	3.67e+04	2.034	0.042	2728.027	1.46e+05
zipcode_98022	8.621e+04	2.03e+04	4.253	0.000	4.65e+04	1.26e+05
zipcode_98023	-5.151e+04	1.26e+04	-4.093	0.000	-7.62e+04	-2.68e+04
zipcode_98024	1.806e+05	3.27e+04	5.524	0.000	1.17e+05	2.45e+05
zipcode_98027	1.718e+05	2.23e+04	7.706	0.000	1.28e+05	2.16e+05
zipcode_98028	6.885e+04	3.28e+04	2.097	0.036	4507.689	1.33e+05
zipcode_98029	2.212e+05	2.55e+04	8.683	0.000	1.71e+05	2.71e+05
zipcode_98030	6440.1953	1.5e+04	0.428	0.669	-2.31e+04	3.59e+04
zipcode_98031	1.687e+04	1.57e+04	1.076	0.282	-1.39e+04	4.76e+04
zipcode_98032	9181.2935	1.81e+04	0.507	0.612	-2.63e+04	4.47e+04
zipcode_98033	3.43e+05	2.81e+04	12.190	0.000	2.88e+05	3.98e+05
zipcode_98034	1.685e+05	3.02e+04	5.583	0.000	1.09e+05	2.28e+05
zipcode_98038	5.275e+04	1.69e+04	3.115	0.002	1.96e+04	8.59e+04
zipcode_98039	1.275e+06	3.35e+04	38.079	0.000	1.21e+06	1.34e+06
zipcode_98040	5.198e+05	2.19e+04	23.764	0.000	4.77e+05	5.63e+05
zipcode_98042	2.321e+04	1.44e+04	1.615	0.106	-4962.935	5.14e+04
zipcode_98045	1.575e+05	3.13e+04	5.039	0.000	9.62e+04	2.19e+05
zipcode_98052	1.962e+05	2.88e+04	6.825	0.000	1.4e+05	2.53e+05
zipcode_98053	1.611e+05	3.08e+04	5.224	0.000	1.01e+05	2.21e+05
zipcode_98055	4.754e+04	1.74e+04	2.729	0.006	1.34e+04	8.17e+04
zipcode_98056	9.953e+04	1.89e+04	5.274	0.000	6.25e+04	1.37e+05
zipcode_98058	3.033e+04	1.65e+04	1.841	0.066	-1957.843	6.26e+04
zipcode_98059	7.367e+04	1.86e+04	3.969	0.000	3.73e+04	1.1e+05
zipcode_98065	1.18e+05	2.88e+04	4.098	0.000	6.15e+04	1.74e+05
zipcode_98070	-1.88e+04	2.17e+04	-0.867	0.386	-6.13e+04	2.37e+04
zipcode_98072	1.063e+05	3.36e+04	3.160	0.002	4.03e+04	1.72e+05
zipcode_98074	1.576e+05	2.72e+04	5.785	0.000	1.04e+05	2.11e+05
zipcode_98075	1.604e+05	2.62e+04	6.116	0.000	1.09e+05	2.12e+05
zipcode_98077	7.644e+04	3.5e+04	2.185	0.029	7873.688	1.45e+05
zipcode_98092	-2.541e+04	1.37e+04	-1.855	0.064	-5.23e+04	1439.737
zipcode_98102	5.076e+05	2.9e+04	17.532	0.000	4.51e+05	5.64e+05
zipcode_98103	3.306e+05	2.71e+04	12.201	0.000	2.78e+05	3.84e+05

zipcode_98105	4.71e+05	2.78e+04	16.967	0.000	4.17e+05	5.25e+05
zipcode_98106	1.245e+05	2.02e+04	6.177	0.000	8.5e+04	1.64e+05
zipcode_98107	3.323e+05	2.8e+04	11.882	0.000	2.77e+05	3.87e+05
zipcode_98108	1.132e+05	2.22e+04	5.099	0.000	6.97e+04	1.57e+05
zipcode_98109	4.99e+05	2.88e+04	17.319	0.000	4.43e+05	5.55e+05
zipcode_98112	6.152e+05	2.55e+04	24.168	0.000	5.65e+05	6.65e+05
zipcode_98115	3.155e+05	2.76e+04	11.436	0.000	2.61e+05	3.7e+05
zipcode_98116	3.002e+05	2.24e+04	13.379	0.000	2.56e+05	3.44e+05
zipcode_98117	2.948e+05	2.79e+04	10.552	0.000	2.4e+05	3.5e+05
zipcode_98118	1.769e+05	1.96e+04	9.036	0.000	1.39e+05	2.15e+05
zipcode_98119	4.967e+05	2.72e+04	18.259	0.000	4.43e+05	5.5e+05
zipcode_98122	3.457e+05	2.42e+04	14.279	0.000	2.98e+05	3.93e+05
zipcode_98125	1.726e+05	2.99e+04	5.780	0.000	1.14e+05	2.31e+05
zipcode_98126	1.959e+05	2.06e+04	9.494	0.000	1.55e+05	2.36e+05
zipcode_98133	1.233e+05	3.09e+04	3.996	0.000	6.28e+04	1.84e+05
zipcode_98136	2.492e+05	2.12e+04	11.770	0.000	2.08e+05	2.91e+05
zipcode_98144	2.904e+05	2.25e+04	12.881	0.000	2.46e+05	3.35e+05
zipcode_98146	1.078e+05	1.89e+04	5.692	0.000	7.07e+04	1.45e+05
zipcode_98148	4.939e+04	2.59e+04	1.907	0.057	-1381.603	1e+05
zipcode_98155	1.051e+05	3.21e+04	3.275	0.001	4.22e+04	1.68e+05
zipcode_98166	6.379e+04	1.73e+04	3.687	0.000	2.99e+04	9.77e+04
zipcode_98168	6.116e+04	1.83e+04	3.341	0.001	2.53e+04	9.7e+04
zipcode_98177	1.959e+05	3.22e+04	6.090	0.000	1.33e+05	2.59e+05
zipcode_98178	4.907e+04	1.89e+04	2.602	0.009	1.21e+04	8.6e+04
zipcode_98188	3.065e+04	1.95e+04	1.571	0.116	-7588.055	6.89e+04
zipcode_98198	1.591e+04	1.47e+04	1.079	0.281	-1.3e+04	4.48e+04
zipcode_98199	3.709e+05	2.65e+04	13.987	0.000	3.19e+05	4.23e+05


```
=====
Omnibus:                20092.654    Durbin-Watson:                1.985
Prob(Omnibus):           0.000    Jarque-Bera (JB):            3585946.415
Skew:                    4.107    Prob(JB):                     0.00
Kurtosis:                66.270    Cond. No.                     2.47e+08
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.47e+08. This might indicate that there are strong multicollinearity or other numerical problems.

"""



6.2 Removing Outliers to Fulfill Assumptions of Multiple Regressions

We have successfully run our baseline model, and our R2 value isn't too bad! However, we can see from the QQ plot and homoscedasticity plot that we are not fulfilling the assumptions of multiple regression.

We will try to address this issue by removing outliers that lie 1.5 times the IQR below the first quartile and 1.5 times the IQR above the third quartile.

```
[37]: # Create function to remove outliers.
def find_outliers_IQR(data):
    """This function was written by James Irving during study group.

    Detects outliers using the 1.5*IQR thresholds.
    Returns a boolean Series where True=outlier"""
    res = data.describe()
    q1 = res['25%']
    q3 = res['75%']
    thresh = 1.5*(q3-q1)
    idx_outliers =(data < (q1-thresh)) | (data > (q3+thresh))
    return idx_outliers
```

In the 'Explore' section, we saw that we have many outliers several columns. We will proceed to remove outliers from those columns that have extreme outliers, based on our boxplot visualizations.

```
[38]: # Create list of columns to remove outliers from
cols_outlier = ['price', 'bedrooms', 'bathrooms', 'sqft_lot', 'sqft_above', \
                'sqft_basement', 'sqft_living15']
df_outliers = df_model.copy()
```

```
# Remove outliers for specified columns
for col in cols_outlier:
    df_outliers = df_outliers[~find_outliers_IQR(df_outliers[col])]
```

```
[39]: df_outliers
```

```
[39]:
```

	price	bedrooms	bathrooms	sqft_lot	floors	waterfront	grade	\
0	221900.0	3	1.00	5650	1.0	0.0	7	
1	538000.0	3	2.25	7242	2.0	0.0	7	
2	180000.0	2	1.00	10000	1.0	0.0	6	
3	604000.0	4	3.00	5000	1.0	0.0	7	
4	510000.0	3	2.00	8080	1.0	0.0	8	
...	
21592	360000.0	3	2.50	1131	3.0	0.0	8	
21593	400000.0	4	2.50	5813	2.0	0.0	8	
21594	402101.0	2	0.75	1350	2.0	0.0	7	
21595	400000.0	3	2.50	2388	2.0	0.0	8	
21596	325000.0	2	0.75	1076	2.0	0.0	7	

	sqft_above	sqft_basement	lat	...	zipcode_98146	zipcode_98148	\
0	1180	0.0	47.5112	...	0.0	0.0	
1	2170	400.0	47.7210	...	0.0	0.0	
2	770	0.0	47.7379	...	0.0	0.0	
3	1050	910.0	47.5208	...	0.0	0.0	
4	1680	0.0	47.6168	...	0.0	0.0	
...	
21592	1530	0.0	47.6993	...	0.0	0.0	
21593	2310	0.0	47.5107	...	1.0	0.0	
21594	1020	0.0	47.5944	...	0.0	0.0	
21595	1600	0.0	47.5345	...	0.0	0.0	
21596	1020	0.0	47.5941	...	0.0	0.0	

	zipcode_98155	zipcode_98166	zipcode_98168	zipcode_98177	\
0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	
...	
21592	0.0	0.0	0.0	0.0	
21593	0.0	0.0	0.0	0.0	
21594	0.0	0.0	0.0	0.0	
21595	0.0	0.0	0.0	0.0	
21596	0.0	0.0	0.0	0.0	

	zipcode_98178	zipcode_98188	zipcode_98198	zipcode_98199
0	1.0	0.0	0.0	0.0

1	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0
...
21592	0.0	0.0	0.0	0.0
21593	0.0	0.0	0.0	0.0
21594	0.0	0.0	0.0	0.0
21595	0.0	0.0	0.0	0.0
21596	0.0	0.0	0.0	0.0

[16358 rows x 81 columns]

```
[40]: # Run regression model on our dataset where outliers are removed.
model_combined(df_outliers);
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                OLS Regression Results
=====
Dep. Variable:                  price    R-squared:                0.807
Model:                            OLS    Adj. R-squared:           0.806
Method:                 Least Squares    F-statistic:                848.1
Date:                Thu, 22 Apr 2021    Prob (F-statistic):          0.00
Time:                22:29:31    Log-Likelihood:            -2.0856e+05
No. Observations:          16358    AIC:                        4.173e+05
Df Residuals:              16277    BIC:                        4.179e+05
Df Model:                   80
Covariance Type:            nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-2.098e+06	4.45e+06	-0.471	0.638	-1.08e+07	6.63e+06
bedrooms	-2850.0552	1115.288	-2.555	0.011	-5036.141	-663.969
bathrooms	6959.9481	1611.666	4.318	0.000	3800.906	1.01e+04
sqft_lot	2.8752	0.276	10.421	0.000	2.334	3.416
floors	-2.75e+04	1951.375	-14.092	0.000	-3.13e+04	-2.37e+04
waterfront	3.382e+05	1.87e+04	18.087	0.000	3.02e+05	3.75e+05
grade	3.475e+04	1149.769	30.221	0.000	3.25e+04	3.7e+04
sqft_above	130.3372	2.348	55.502	0.000	125.734	134.940
sqft_basement	90.9682	2.638	34.480	0.000	85.797	96.140
lat	-4.173e+04	4.19e+04	-0.995	0.320	-1.24e+05	4.05e+04
long	-3.107e+04	3.37e+04	-0.922	0.356	-9.71e+04	3.5e+04
sqft_living15	34.8750	2.108	16.545	0.000	30.743	39.007
zipcode_98002	3.279e+04	8193.908	4.001	0.000	1.67e+04	4.88e+04
zipcode_98003	5708.5590	7433.852	0.768	0.443	-8862.607	2.03e+04
zipcode_98004	5.463e+05	1.58e+04	34.605	0.000	5.15e+05	5.77e+05

zipcode_98005	3.521e+05	1.62e+04	21.728	0.000	3.2e+05	3.84e+05
zipcode_98006	2.942e+05	1.34e+04	21.879	0.000	2.68e+05	3.21e+05
zipcode_98007	2.793e+05	1.64e+04	17.022	0.000	2.47e+05	3.11e+05
zipcode_98008	2.716e+05	1.59e+04	17.079	0.000	2.4e+05	3.03e+05
zipcode_98010	1.052e+05	1.61e+04	6.535	0.000	7.37e+04	1.37e+05
zipcode_98011	1.656e+05	2.06e+04	8.024	0.000	1.25e+05	2.06e+05
zipcode_98014	1.37e+05	2.61e+04	5.255	0.000	8.59e+04	1.88e+05
zipcode_98019	1.15e+05	2.29e+04	5.023	0.000	7.01e+04	1.6e+05
zipcode_98022	3.168e+04	1.33e+04	2.382	0.017	5610.093	5.78e+04
zipcode_98023	-1.431e+04	7174.382	-1.995	0.046	-2.84e+04	-247.789
zipcode_98024	1.662e+05	2.45e+04	6.796	0.000	1.18e+05	2.14e+05
zipcode_98027	2.528e+05	1.46e+04	17.348	0.000	2.24e+05	2.81e+05
zipcode_98028	1.532e+05	2.01e+04	7.620	0.000	1.14e+05	1.93e+05
zipcode_98029	2.584e+05	1.6e+04	16.180	0.000	2.27e+05	2.9e+05
zipcode_98030	1.107e+04	8322.238	1.330	0.184	-5244.964	2.74e+04
zipcode_98031	2.545e+04	8842.824	2.878	0.004	8117.202	4.28e+04
zipcode_98032	1.577e+04	9774.747	1.614	0.107	-3384.816	3.49e+04
zipcode_98033	3.44e+05	1.74e+04	19.733	0.000	3.1e+05	3.78e+05
zipcode_98034	2.116e+05	1.86e+04	11.375	0.000	1.75e+05	2.48e+05
zipcode_98038	4.761e+04	1.05e+04	4.525	0.000	2.7e+04	6.82e+04
zipcode_98039	6.678e+05	3.71e+04	18.020	0.000	5.95e+05	7.4e+05
zipcode_98040	4.52e+05	1.42e+04	31.871	0.000	4.24e+05	4.8e+05
zipcode_98042	2.367e+04	8705.330	2.719	0.007	6610.572	4.07e+04
zipcode_98045	1.206e+05	2.06e+04	5.851	0.000	8.02e+04	1.61e+05
zipcode_98052	2.762e+05	1.77e+04	15.598	0.000	2.41e+05	3.11e+05
zipcode_98053	2.734e+05	2.02e+04	13.520	0.000	2.34e+05	3.13e+05
zipcode_98055	6.081e+04	1e+04	6.051	0.000	4.11e+04	8.05e+04
zipcode_98056	1.314e+05	1.12e+04	11.688	0.000	1.09e+05	1.53e+05
zipcode_98058	5.115e+04	9796.055	5.221	0.000	3.19e+04	7.03e+04
zipcode_98059	1.02e+05	1.12e+04	9.127	0.000	8.01e+04	1.24e+05
zipcode_98065	1.585e+05	1.86e+04	8.521	0.000	1.22e+05	1.95e+05
zipcode_98070	8.554e+04	1.89e+04	4.521	0.000	4.85e+04	1.23e+05
zipcode_98072	1.76e+05	2.12e+04	8.312	0.000	1.35e+05	2.18e+05
zipcode_98074	2.267e+05	1.72e+04	13.149	0.000	1.93e+05	2.6e+05
zipcode_98075	2.514e+05	1.72e+04	14.627	0.000	2.18e+05	2.85e+05
zipcode_98077	1.773e+05	2.61e+04	6.798	0.000	1.26e+05	2.28e+05
zipcode_98092	-1.667e+04	7912.858	-2.107	0.035	-3.22e+04	-1159.382
zipcode_98102	4.591e+05	1.73e+04	26.564	0.000	4.25e+05	4.93e+05
zipcode_98103	3.806e+05	1.66e+04	22.896	0.000	3.48e+05	4.13e+05
zipcode_98105	4.348e+05	1.7e+04	25.507	0.000	4.01e+05	4.68e+05
zipcode_98106	1.511e+05	1.2e+04	12.608	0.000	1.28e+05	1.75e+05
zipcode_98107	3.775e+05	1.7e+04	22.229	0.000	3.44e+05	4.11e+05
zipcode_98108	1.532e+05	1.3e+04	11.825	0.000	1.28e+05	1.79e+05
zipcode_98109	4.707e+05	1.74e+04	27.075	0.000	4.37e+05	5.05e+05
zipcode_98112	4.866e+05	1.58e+04	30.810	0.000	4.56e+05	5.18e+05
zipcode_98115	3.67e+05	1.69e+04	21.701	0.000	3.34e+05	4e+05
zipcode_98116	3.521e+05	1.35e+04	25.985	0.000	3.25e+05	3.79e+05
zipcode_98117	3.64e+05	1.72e+04	21.192	0.000	3.3e+05	3.98e+05

zipcode_98118	2.024e+05	1.17e+04	17.309	0.000	1.8e+05	2.25e+05
zipcode_98119	4.688e+05	1.65e+04	28.358	0.000	4.36e+05	5.01e+05
zipcode_98122	3.629e+05	1.46e+04	24.918	0.000	3.34e+05	3.91e+05
zipcode_98125	2.325e+05	1.83e+04	12.722	0.000	1.97e+05	2.68e+05
zipcode_98126	2.406e+05	1.23e+04	19.517	0.000	2.16e+05	2.65e+05
zipcode_98133	1.894e+05	1.89e+04	9.994	0.000	1.52e+05	2.27e+05
zipcode_98136	3.027e+05	1.26e+04	24.041	0.000	2.78e+05	3.27e+05
zipcode_98144	2.933e+05	1.36e+04	21.592	0.000	2.67e+05	3.2e+05
zipcode_98146	1.346e+05	1.11e+04	12.104	0.000	1.13e+05	1.56e+05
zipcode_98148	6.492e+04	1.37e+04	4.744	0.000	3.81e+04	9.17e+04
zipcode_98155	1.728e+05	1.97e+04	8.771	0.000	1.34e+05	2.11e+05
zipcode_98166	1.211e+05	1.03e+04	11.763	0.000	1.01e+05	1.41e+05
zipcode_98168	6.63e+04	1.07e+04	6.195	0.000	4.53e+04	8.73e+04
zipcode_98177	2.393e+05	1.98e+04	12.090	0.000	2.01e+05	2.78e+05
zipcode_98178	8.289e+04	1.09e+04	7.593	0.000	6.15e+04	1.04e+05
zipcode_98188	5.228e+04	1.09e+04	4.790	0.000	3.09e+04	7.37e+04
zipcode_98198	4.982e+04	8274.755	6.020	0.000	3.36e+04	6.6e+04
zipcode_98199	4.024e+05	1.63e+04	24.658	0.000	3.7e+05	4.34e+05


```
=====
Omnibus:                1839.987    Durbin-Watson:           2.005
Prob(Omnibus):           0.000    Jarque-Bera (JB):       6354.147
Skew:                    0.558    Prob(JB):               0.00
Kurtosis:                5.842    Cond. No.               5.56e+07
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.56e+07. This might indicate that there are strong multicollinearity or other numerical problems.

"""



Great! We can see that although they are not quite perfect, our QQ plot and homoscedasticity plot look much better. We can see that our R2 value has gone up a bit as well.

Now we want to move on to addressing the nonsignificant P-values in our model. Since a nonsignificant P-value indicates that our model would be no different than when the respective coefficient is 0, we will go ahead and remove those variables from our model.

```
[41]: # Create new dataframe after removing outliers
df_pvalues = df_outliers.drop(['lat', 'long'], axis=1)
```

```
[42]: # Run regression on dataframe after removing non-significant variables
model_unscaled, fig_unscaled, ax_unscaled = model_combined(df_pvalues)
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""

                        OLS Regression Results
=====
Dep. Variable:          price      R-squared:                0.806
Model:                  OLS      Adj. R-squared:             0.806
Method:                 Least Squares    F-statistic:          869.8
Date:                  Thu, 22 Apr 2021    Prob (F-statistic):      0.00
Time:                  22:29:31    Log-Likelihood:         -2.0856e+05
No. Observations:      16358    AIC:                    4.173e+05
Df Residuals:          16279    BIC:                    4.179e+05
Df Model:               78
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-2.73e+05	8978.380	-30.405	0.000	-2.91e+05	-2.55e+05
bedrooms	-2849.0091	1115.274	-2.555	0.011	-5035.068	-662.950
bathrooms	6955.4124	1611.577	4.316	0.000	3796.545	1.01e+04
sqft_lot	2.8718	0.276	10.409	0.000	2.331	3.413
floors	-2.748e+04	1951.119	-14.082	0.000	-3.13e+04	-2.37e+04
waterfront	3.385e+05	1.87e+04	18.106	0.000	3.02e+05	3.75e+05
grade	3.476e+04	1149.100	30.250	0.000	3.25e+04	3.7e+04
sqft_above	130.3278	2.348	55.500	0.000	125.725	134.931
sqft_basement	90.9779	2.638	34.486	0.000	85.807	96.149
sqft_living15	34.8325	2.107	16.528	0.000	30.702	38.963
zipcode_98002	3.111e+04	7975.317	3.901	0.000	1.55e+04	4.67e+04
zipcode_98003	6760.4627	7310.504	0.925	0.355	-7568.927	2.11e+04
zipcode_98004	5.316e+05	9106.136	58.382	0.000	5.14e+05	5.49e+05
zipcode_98005	3.368e+05	9855.624	34.175	0.000	3.17e+05	3.56e+05
zipcode_98006	2.802e+05	7147.222	39.199	0.000	2.66e+05	2.94e+05
zipcode_98007	2.628e+05	9275.845	28.334	0.000	2.45e+05	2.81e+05
zipcode_98008	2.541e+05	7420.469	34.249	0.000	2.4e+05	2.69e+05

zipcode_98010	9.676e+04	1.36e+04	7.131	0.000	7.02e+04	1.23e+05
zipcode_98011	1.45e+05	8316.130	17.434	0.000	1.29e+05	1.61e+05
zipcode_98014	1.08e+05	1.39e+04	7.790	0.000	8.08e+04	1.35e+05
zipcode_98019	8.778e+04	8780.990	9.997	0.000	7.06e+04	1.05e+05
zipcode_98022	2.768e+04	8721.026	3.174	0.002	1.06e+04	4.48e+04
zipcode_98023	-1.125e+04	6426.804	-1.751	0.080	-2.39e+04	1342.806
zipcode_98024	1.443e+05	1.78e+04	8.107	0.000	1.09e+05	1.79e+05
zipcode_98027	2.366e+05	7721.041	30.640	0.000	2.21e+05	2.52e+05
zipcode_98028	1.339e+05	7433.627	18.012	0.000	1.19e+05	1.48e+05
zipcode_98029	2.397e+05	7261.295	33.009	0.000	2.25e+05	2.54e+05
zipcode_98030	6138.2120	7420.942	0.827	0.408	-8407.648	2.07e+04
zipcode_98031	1.904e+04	7352.505	2.590	0.010	4630.856	3.35e+04
zipcode_98032	1.343e+04	9420.742	1.426	0.154	-5032.423	3.19e+04
zipcode_98033	3.259e+05	6946.511	46.914	0.000	3.12e+05	3.4e+05
zipcode_98034	1.924e+05	6329.075	30.403	0.000	1.8e+05	2.05e+05
zipcode_98038	3.808e+04	6374.554	5.974	0.000	2.56e+04	5.06e+04
zipcode_98039	6.533e+05	3.45e+04	18.938	0.000	5.86e+05	7.21e+05
zipcode_98040	4.401e+05	9275.468	47.452	0.000	4.22e+05	4.58e+05
zipcode_98042	1.654e+04	6449.312	2.565	0.010	3897.958	2.92e+04
zipcode_98045	9.826e+04	8771.005	11.203	0.000	8.11e+04	1.15e+05
zipcode_98052	2.562e+05	6441.635	39.777	0.000	2.44e+05	2.69e+05
zipcode_98053	2.497e+05	7924.596	31.508	0.000	2.34e+05	2.65e+05
zipcode_98055	5.255e+04	7470.196	7.035	0.000	3.79e+04	6.72e+04
zipcode_98056	1.204e+05	6723.038	17.906	0.000	1.07e+05	1.34e+05
zipcode_98058	4.178e+04	6661.999	6.272	0.000	2.87e+04	5.48e+04
zipcode_98059	9.057e+04	6761.772	13.394	0.000	7.73e+04	1.04e+05
zipcode_98065	1.368e+05	7640.870	17.898	0.000	1.22e+05	1.52e+05
zipcode_98070	8.72e+04	1.74e+04	5.009	0.000	5.31e+04	1.21e+05
zipcode_98072	1.537e+05	8855.588	17.355	0.000	1.36e+05	1.71e+05
zipcode_98074	2.065e+05	7135.626	28.933	0.000	1.92e+05	2.2e+05
zipcode_98075	2.322e+05	8918.569	26.034	0.000	2.15e+05	2.5e+05
zipcode_98077	1.529e+05	1.69e+04	9.066	0.000	1.2e+05	1.86e+05
zipcode_98092	-1.899e+04	7377.054	-2.575	0.010	-3.35e+04	-4532.805
zipcode_98102	4.472e+05	1.07e+04	41.757	0.000	4.26e+05	4.68e+05
zipcode_98103	3.677e+05	6345.370	57.943	0.000	3.55e+05	3.8e+05
zipcode_98105	4.209e+05	8354.953	50.372	0.000	4.04e+05	4.37e+05
zipcode_98106	1.443e+05	7015.786	20.573	0.000	1.31e+05	1.58e+05
zipcode_98107	3.657e+05	7473.143	48.939	0.000	3.51e+05	3.8e+05
zipcode_98108	1.443e+05	8224.265	17.550	0.000	1.28e+05	1.6e+05
zipcode_98109	4.596e+05	1.07e+04	42.874	0.000	4.39e+05	4.81e+05
zipcode_98112	4.742e+05	8524.249	55.627	0.000	4.57e+05	4.91e+05
zipcode_98115	3.523e+05	6321.102	55.740	0.000	3.4e+05	3.65e+05
zipcode_98116	3.449e+05	7158.996	48.178	0.000	3.31e+05	3.59e+05
zipcode_98117	3.517e+05	6393.125	55.010	0.000	3.39e+05	3.64e+05
zipcode_98118	1.929e+05	6445.325	29.926	0.000	1.8e+05	2.06e+05
zipcode_98119	4.58e+05	8775.961	52.187	0.000	4.41e+05	4.75e+05
zipcode_98122	3.514e+05	7434.646	47.260	0.000	3.37e+05	3.66e+05
zipcode_98125	2.165e+05	6684.149	32.394	0.000	2.03e+05	2.3e+05

zipcode_98126	2.341e+05	6939.693	33.740	0.000	2.21e+05	2.48e+05
zipcode_98133	1.739e+05	6376.183	27.274	0.000	1.61e+05	1.86e+05
zipcode_98136	2.968e+05	7511.155	39.511	0.000	2.82e+05	3.11e+05
zipcode_98144	2.827e+05	7165.040	39.455	0.000	2.69e+05	2.97e+05
zipcode_98146	1.294e+05	7342.058	17.621	0.000	1.15e+05	1.44e+05
zipcode_98148	6.152e+04	1.25e+04	4.912	0.000	3.7e+04	8.61e+04
zipcode_98155	1.554e+05	6586.017	23.588	0.000	1.42e+05	1.68e+05
zipcode_98166	1.174e+05	7950.541	14.771	0.000	1.02e+05	1.33e+05
zipcode_98168	6.001e+04	7585.424	7.912	0.000	4.51e+04	7.49e+04
zipcode_98177	2.244e+05	7994.410	28.071	0.000	2.09e+05	2.4e+05
zipcode_98178	7.431e+04	7442.250	9.985	0.000	5.97e+04	8.89e+04
zipcode_98188	4.692e+04	9296.686	5.047	0.000	2.87e+04	6.51e+04
zipcode_98198	4.784e+04	7434.966	6.435	0.000	3.33e+04	6.24e+04
zipcode_98199	3.922e+05	7476.812	52.459	0.000	3.78e+05	4.07e+05

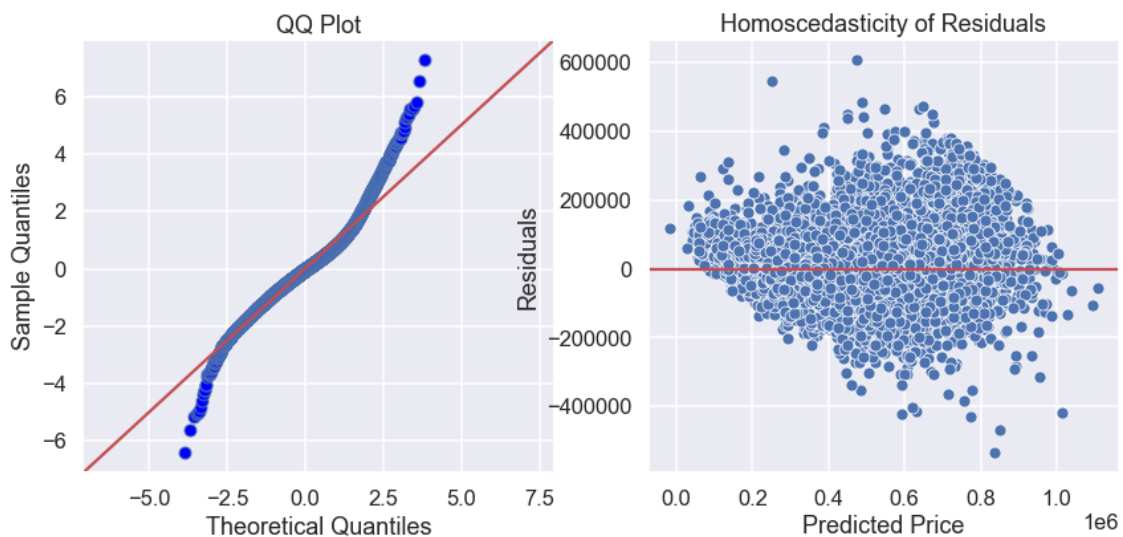
Omnibus:	1844.788	Durbin-Watson:	2.004
Prob(Omnibus):	0.000	Jarque-Bera (JB):	6380.863
Skew:	0.559	Prob(JB):	0.00
Kurtosis:	5.848	Cond. No.	5.37e+05

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.37e+05. This might indicate that there are strong multicollinearity or other numerical problems.

"""



7 iNTERPRET

Now that we have our final model with outliers removed and only significant P-values included, all that's left in our analysis is to scale our model coefficients to determine which coefficients have the largest effect on the variability of housing price. Since there are multiple coefficients for zipcode, we will examine which of the other variables have high coefficients.

We should also note that zipcode, as well as some other variables are ones that we cannot control, and therefore will not be appropriate variables to provide recommendations for changing. However, we will still include those variables as part of our model, as long as they have a high enough coefficient to indicate that they are valid predictors for the value of a house.

7.1 Scaling the Dataset

```
[43]: # Create copy of final dataset to scale
df_unscaled = df_pvalues.copy()
```

```
[44]: # Create list of columns except for zipcode
numeric_cols = [col for col in df_unscaled.columns if \
                 col.startswith('zipcode')==False]
numeric_cols
```

```
[44]: ['price',
       'bedrooms',
       'bathrooms',
       'sqft_lot',
       'floors',
       'waterfront',
       'grade',
       'sqft_above',
       'sqft_basement',
       'sqft_living15']
```

```
[45]: # Create scaler object
scaler = StandardScaler()
scaler
```

```
[45]: StandardScaler()
```

```
[46]: # Scale our dataset used to form our final model
df_scaled = df_unscaled.copy()
df_scaled[numeric_cols] = scaler.fit_transform(df_scaled[numeric_cols])
df_scaled.describe().round(2)
```

```
[46]:
```

	price	bedrooms	bathrooms	sqft_lot	floors	waterfront	\
count	16358.00	16358.00	16358.00	16358.00	16358.00	16358.00	
mean	-0.00	-0.00	0.00	0.00	0.00	-0.00	
std	1.00	1.00	1.00	1.00	1.00	1.00	

min	-1.95	-1.62	-2.24	-1.93	-0.85	-0.04
25%	-0.77	-0.32	-0.72	-0.69	-0.85	-0.04
50%	-0.18	-0.32	0.04	0.00	-0.85	-0.04
75%	0.59	0.98	0.80	0.58	0.98	-0.04
max	3.55	2.28	3.08	3.28	3.74	27.89

	grade	sqft_above	sqft_basement	sqft_living15	...	zipcode_98146	\
count	16358.00	16358.00	16358.00	16358.00	...	16358.00	
mean	-0.00	0.00	-0.00	-0.00	...	0.01	
std	1.00	1.00	1.00	1.00	...	0.12	
min	-3.83	-1.96	-0.67	-2.71	...	0.00	
25%	-0.44	-0.75	-0.67	-0.74	...	0.00	
50%	-0.44	-0.24	-0.67	-0.17	...	0.00	
75%	0.69	0.59	0.68	0.63	...	0.00	
max	4.08	2.91	2.91	2.80	...	1.00	

	zipcode_98148	zipcode_98155	zipcode_98166	zipcode_98168	\
count	16358.00	16358.00	16358.00	16358.00	
mean	0.00	0.02	0.01	0.01	
std	0.06	0.15	0.11	0.11	
min	0.00	0.00	0.00	0.00	
25%	0.00	0.00	0.00	0.00	
50%	0.00	0.00	0.00	0.00	
75%	0.00	0.00	0.00	0.00	
max	1.00	1.00	1.00	1.00	

	zipcode_98177	zipcode_98178	zipcode_98188	zipcode_98198	\
count	16358.00	16358.00	16358.00	16358.00	
mean	0.01	0.01	0.01	0.01	
std	0.10	0.12	0.08	0.12	
min	0.00	0.00	0.00	0.00	
25%	0.00	0.00	0.00	0.00	
50%	0.00	0.00	0.00	0.00	
75%	0.00	0.00	0.00	0.00	
max	1.00	1.00	1.00	1.00	

	zipcode_98199
count	16358.00
mean	0.01
std	0.12
min	0.00
25%	0.00
50%	0.00
75%	0.00
max	1.00

[8 rows x 79 columns]

7.2 Creating a Scaled Model

```
[47]: # Run regression model on scaled data
model_scaled, fig_scaled, ax_scaled = model_combined(df_scaled)
```

```
<class 'statsmodels.iolib.summary.Summary'>
```

```
"""
```

OLS Regression Results

```
=====
Dep. Variable:          price    R-squared:                0.806
Model:                  OLS      Adj. R-squared:            0.806
Method:                 Least Squares    F-statistic:          869.8
Date:                   Thu, 22 Apr 2021    Prob (F-statistic):      0.00
Time:                   22:29:33    Log-Likelihood:         -9777.5
No. Observations:      16358    AIC:                   1.971e+04
Df Residuals:          16279    BIC:                   2.032e+04
Df Model:               78
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-1.0046	0.027	-37.858	0.000	-1.057	-0.953
bedrooms	-0.0116	0.005	-2.555	0.011	-0.020	-0.003
bathrooms	0.0242	0.006	4.316	0.000	0.013	0.035
sqft_lot	0.0514	0.005	10.409	0.000	0.042	0.061
floors	-0.0790	0.006	-14.082	0.000	-0.090	-0.068
waterfront	0.0640	0.004	18.106	0.000	0.057	0.071
grade	0.1623	0.005	30.250	0.000	0.152	0.173
sqft_above	0.3970	0.007	55.500	0.000	0.383	0.411
sqft_basement	0.1700	0.005	34.486	0.000	0.160	0.180
sqft_living15	0.0908	0.005	16.528	0.000	0.080	0.102
zipcode_98002	0.1642	0.042	3.901	0.000	0.082	0.247
zipcode_98003	0.0357	0.039	0.925	0.355	-0.040	0.111
zipcode_98004	2.8061	0.048	58.382	0.000	2.712	2.900
zipcode_98005	1.7778	0.052	34.175	0.000	1.676	1.880
zipcode_98006	1.4788	0.038	39.199	0.000	1.405	1.553
zipcode_98007	1.3872	0.049	28.334	0.000	1.291	1.483
zipcode_98008	1.3414	0.039	34.249	0.000	1.265	1.418
zipcode_98010	0.5107	0.072	7.131	0.000	0.370	0.651
zipcode_98011	0.7653	0.044	17.434	0.000	0.679	0.851
zipcode_98014	0.5701	0.073	7.790	0.000	0.427	0.714
zipcode_98019	0.4633	0.046	9.997	0.000	0.372	0.554
zipcode_98022	0.1461	0.046	3.174	0.002	0.056	0.236
zipcode_98023	-0.0594	0.034	-1.751	0.080	-0.126	0.007
zipcode_98024	0.7616	0.094	8.107	0.000	0.577	0.946
zipcode_98027	1.2487	0.041	30.640	0.000	1.169	1.329
zipcode_98028	0.7067	0.039	18.012	0.000	0.630	0.784
zipcode_98029	1.2652	0.038	33.009	0.000	1.190	1.340


zipcode_98030	0.0324	0.039	0.827	0.408	-0.044	0.109
zipcode_98031	0.1005	0.039	2.590	0.010	0.024	0.177
zipcode_98032	0.0709	0.050	1.426	0.154	-0.027	0.168
zipcode_98033	1.7202	0.037	46.914	0.000	1.648	1.792
zipcode_98034	1.0157	0.033	30.403	0.000	0.950	1.081
zipcode_98038	0.2010	0.034	5.974	0.000	0.135	0.267
zipcode_98039	3.4485	0.182	18.938	0.000	3.092	3.805
zipcode_98040	2.3232	0.049	47.452	0.000	2.227	2.419
zipcode_98042	0.0873	0.034	2.565	0.010	0.021	0.154
zipcode_98045	0.5187	0.046	11.203	0.000	0.428	0.609
zipcode_98052	1.3525	0.034	39.777	0.000	1.286	1.419
zipcode_98053	1.3179	0.042	31.508	0.000	1.236	1.400
zipcode_98055	0.2774	0.039	7.035	0.000	0.200	0.355
zipcode_98056	0.6354	0.035	17.906	0.000	0.566	0.705
zipcode_98058	0.2205	0.035	6.272	0.000	0.152	0.289
zipcode_98059	0.4780	0.036	13.394	0.000	0.408	0.548
zipcode_98065	0.7218	0.040	17.898	0.000	0.643	0.801
zipcode_98070	0.4602	0.092	5.009	0.000	0.280	0.640
zipcode_98072	0.8112	0.047	17.355	0.000	0.720	0.903
zipcode_98074	1.0897	0.038	28.933	0.000	1.016	1.164
zipcode_98075	1.2255	0.047	26.034	0.000	1.133	1.318
zipcode_98077	0.8072	0.089	9.066	0.000	0.633	0.982
zipcode_98092	-0.1002	0.039	-2.575	0.010	-0.177	-0.024
zipcode_98102	2.3603	0.057	41.757	0.000	2.250	2.471
zipcode_98103	1.9407	0.033	57.943	0.000	1.875	2.006
zipcode_98105	2.2214	0.044	50.372	0.000	2.135	2.308
zipcode_98106	0.7618	0.037	20.573	0.000	0.689	0.834
zipcode_98107	1.9304	0.039	48.939	0.000	1.853	2.008
zipcode_98108	0.7618	0.043	17.550	0.000	0.677	0.847
zipcode_98109	2.4259	0.057	42.874	0.000	2.315	2.537
zipcode_98112	2.5028	0.045	55.627	0.000	2.415	2.591
zipcode_98115	1.8598	0.033	55.740	0.000	1.794	1.925
zipcode_98116	1.8205	0.038	48.178	0.000	1.746	1.895
zipcode_98117	1.8563	0.034	55.010	0.000	1.790	1.922
zipcode_98118	1.0181	0.034	29.926	0.000	0.951	1.085
zipcode_98119	2.4174	0.046	52.187	0.000	2.327	2.508
zipcode_98122	1.8546	0.039	47.260	0.000	1.778	1.932
zipcode_98125	1.1429	0.035	32.394	0.000	1.074	1.212
zipcode_98126	1.2359	0.037	33.740	0.000	1.164	1.308
zipcode_98133	0.9179	0.034	27.274	0.000	0.852	0.984
zipcode_98136	1.5665	0.040	39.511	0.000	1.489	1.644
zipcode_98144	1.4922	0.038	39.455	0.000	1.418	1.566
zipcode_98146	0.6829	0.039	17.621	0.000	0.607	0.759
zipcode_98148	0.3247	0.066	4.912	0.000	0.195	0.454
zipcode_98155	0.8200	0.035	23.588	0.000	0.752	0.888
zipcode_98166	0.6199	0.042	14.771	0.000	0.538	0.702
zipcode_98168	0.3168	0.040	7.912	0.000	0.238	0.395
zipcode_98177	1.1845	0.042	28.071	0.000	1.102	1.267

zipcode_98178	0.3922	0.039	9.985	0.000	0.315	0.469
zipcode_98188	0.2477	0.049	5.047	0.000	0.151	0.344
zipcode_98198	0.2525	0.039	6.435	0.000	0.176	0.329
zipcode_98199	2.0703	0.039	52.459	0.000	1.993	2.148

Omnibus:	1844.788	Durbin-Watson:	2.004
Prob(Omnibus):	0.000	Jarque-Bera (JB):	6380.863
Skew:	0.559	Prob(JB):	0.00
Kurtosis:	5.848	Cond. No.	122.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



7.3 Selecting Variables to Recommend

Now that we have a scaled model, we can pick out the variables with the highest coefficients. This means that we are selecting variables which have the largest impact on the variability of the value of a house.

```
[48]: # Create dataframe of coefficients sorted by highest absolute value
coeffs = model_scaled.params.sort_values().to_frame('coeffs')
coeffs['abs'] = coeffs['coeffs'].abs()
coeffs.sort_values('abs', ascending=False, inplace=True)
coeffs.reset_index(inplace=True)
coeffs[~coeffs['index'].str.startswith('zipcode')]
```

```
[48]:          index      coeffs      abs
33      Intercept -1.004551  1.004551
53      sqft_above  0.397017  0.397017
62  sqft_basement  0.170033  0.170033
64          grade  0.162327  0.162327
68  sqft_living15  0.090830  0.090830
70          floors -0.078977  0.078977
72      waterfront  0.063982  0.063982
74          sqft_lot  0.051372  0.051372
77      bathrooms  0.024169  0.024169
78      bedrooms -0.011559  0.011559
```

We can see that aside from the intercept, our coefficients for 'sqft_above', 'sqft_basement', and 'grade' have the most impact on price. Therefore, we will select those variables to interpret and make recommendations to our stakeholder on.

```
[49]: model_unscaled.summary()
```

```
[49]: <class 'statsmodels.iolib.summary.Summary'>
      """
              OLS Regression Results
=====
Dep. Variable:          price      R-squared:          0.806
Model:                  OLS      Adj. R-squared:       0.806
Method:                 Least Squares      F-statistic:      869.8
Date:                   Thu, 22 Apr 2021      Prob (F-statistic):    0.00
Time:                   22:29:33      Log-Likelihood:      -2.0856e+05
No. Observations:      16358      AIC:                4.173e+05
Df Residuals:          16279      BIC:                4.179e+05
Df Model:               78
Covariance Type:        nonrobust
=====
=
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-
Intercept      -2.73e+05    8978.380    -30.405     0.000    -2.91e+05
-2.55e+05
bedrooms      -2849.0091    1115.274     -2.555     0.011    -5035.068
-662.950
bathrooms       6955.4124    1611.577      4.316     0.000     3796.545
1.01e+04
sqft_lot         2.8718         0.276     10.409     0.000         2.331
3.413
floors        -2.748e+04    1951.119    -14.082     0.000    -3.13e+04
-2.37e+04
waterfront     3.385e+05    1.87e+04     18.106     0.000     3.02e+05
```

3.75e+05					
grade	3.476e+04	1149.100	30.250	0.000	3.25e+04
3.7e+04					
sqft_above	130.3278	2.348	55.500	0.000	125.725
134.931					
sqft_basement	90.9779	2.638	34.486	0.000	85.807
96.149					
sqft_living15	34.8325	2.107	16.528	0.000	30.702
38.963					
zipcode_98002	3.111e+04	7975.317	3.901	0.000	1.55e+04
4.67e+04					
zipcode_98003	6760.4627	7310.504	0.925	0.355	-7568.927
2.11e+04					
zipcode_98004	5.316e+05	9106.136	58.382	0.000	5.14e+05
5.49e+05					
zipcode_98005	3.368e+05	9855.624	34.175	0.000	3.17e+05
3.56e+05					
zipcode_98006	2.802e+05	7147.222	39.199	0.000	2.66e+05
2.94e+05					
zipcode_98007	2.628e+05	9275.845	28.334	0.000	2.45e+05
2.81e+05					
zipcode_98008	2.541e+05	7420.469	34.249	0.000	2.4e+05
2.69e+05					
zipcode_98010	9.676e+04	1.36e+04	7.131	0.000	7.02e+04
1.23e+05					
zipcode_98011	1.45e+05	8316.130	17.434	0.000	1.29e+05
1.61e+05					
zipcode_98014	1.08e+05	1.39e+04	7.790	0.000	8.08e+04
1.35e+05					
zipcode_98019	8.778e+04	8780.990	9.997	0.000	7.06e+04
1.05e+05					
zipcode_98022	2.768e+04	8721.026	3.174	0.002	1.06e+04
4.48e+04					
zipcode_98023	-1.125e+04	6426.804	-1.751	0.080	-2.39e+04
1342.806					
zipcode_98024	1.443e+05	1.78e+04	8.107	0.000	1.09e+05
1.79e+05					
zipcode_98027	2.366e+05	7721.041	30.640	0.000	2.21e+05
2.52e+05					
zipcode_98028	1.339e+05	7433.627	18.012	0.000	1.19e+05
1.48e+05					
zipcode_98029	2.397e+05	7261.295	33.009	0.000	2.25e+05
2.54e+05					
zipcode_98030	6138.2120	7420.942	0.827	0.408	-8407.648
2.07e+04					
zipcode_98031	1.904e+04	7352.505	2.590	0.010	4630.856
3.35e+04					

zipcode_98032	1.343e+04	9420.742	1.426	0.154	-5032.423
3.19e+04					
zipcode_98033	3.259e+05	6946.511	46.914	0.000	3.12e+05
3.4e+05					
zipcode_98034	1.924e+05	6329.075	30.403	0.000	1.8e+05
2.05e+05					
zipcode_98038	3.808e+04	6374.554	5.974	0.000	2.56e+04
5.06e+04					
zipcode_98039	6.533e+05	3.45e+04	18.938	0.000	5.86e+05
7.21e+05					
zipcode_98040	4.401e+05	9275.468	47.452	0.000	4.22e+05
4.58e+05					
zipcode_98042	1.654e+04	6449.312	2.565	0.010	3897.958
2.92e+04					
zipcode_98045	9.826e+04	8771.005	11.203	0.000	8.11e+04
1.15e+05					
zipcode_98052	2.562e+05	6441.635	39.777	0.000	2.44e+05
2.69e+05					
zipcode_98053	2.497e+05	7924.596	31.508	0.000	2.34e+05
2.65e+05					
zipcode_98055	5.255e+04	7470.196	7.035	0.000	3.79e+04
6.72e+04					
zipcode_98056	1.204e+05	6723.038	17.906	0.000	1.07e+05
1.34e+05					
zipcode_98058	4.178e+04	6661.999	6.272	0.000	2.87e+04
5.48e+04					
zipcode_98059	9.057e+04	6761.772	13.394	0.000	7.73e+04
1.04e+05					
zipcode_98065	1.368e+05	7640.870	17.898	0.000	1.22e+05
1.52e+05					
zipcode_98070	8.72e+04	1.74e+04	5.009	0.000	5.31e+04
1.21e+05					
zipcode_98072	1.537e+05	8855.588	17.355	0.000	1.36e+05
1.71e+05					
zipcode_98074	2.065e+05	7135.626	28.933	0.000	1.92e+05
2.2e+05					
zipcode_98075	2.322e+05	8918.569	26.034	0.000	2.15e+05
2.5e+05					
zipcode_98077	1.529e+05	1.69e+04	9.066	0.000	1.2e+05
1.86e+05					
zipcode_98092	-1.899e+04	7377.054	-2.575	0.010	-3.35e+04
-4532.805					
zipcode_98102	4.472e+05	1.07e+04	41.757	0.000	4.26e+05
4.68e+05					
zipcode_98103	3.677e+05	6345.370	57.943	0.000	3.55e+05
3.8e+05					
zipcode_98105	4.209e+05	8354.953	50.372	0.000	4.04e+05

4.37e+05					
zipcode_98106	1.443e+05	7015.786	20.573	0.000	1.31e+05
1.58e+05					
zipcode_98107	3.657e+05	7473.143	48.939	0.000	3.51e+05
3.8e+05					
zipcode_98108	1.443e+05	8224.265	17.550	0.000	1.28e+05
1.6e+05					
zipcode_98109	4.596e+05	1.07e+04	42.874	0.000	4.39e+05
4.81e+05					
zipcode_98112	4.742e+05	8524.249	55.627	0.000	4.57e+05
4.91e+05					
zipcode_98115	3.523e+05	6321.102	55.740	0.000	3.4e+05
3.65e+05					
zipcode_98116	3.449e+05	7158.996	48.178	0.000	3.31e+05
3.59e+05					
zipcode_98117	3.517e+05	6393.125	55.010	0.000	3.39e+05
3.64e+05					
zipcode_98118	1.929e+05	6445.325	29.926	0.000	1.8e+05
2.06e+05					
zipcode_98119	4.58e+05	8775.961	52.187	0.000	4.41e+05
4.75e+05					
zipcode_98122	3.514e+05	7434.646	47.260	0.000	3.37e+05
3.66e+05					
zipcode_98125	2.165e+05	6684.149	32.394	0.000	2.03e+05
2.3e+05					
zipcode_98126	2.341e+05	6939.693	33.740	0.000	2.21e+05
2.48e+05					
zipcode_98133	1.739e+05	6376.183	27.274	0.000	1.61e+05
1.86e+05					
zipcode_98136	2.968e+05	7511.155	39.511	0.000	2.82e+05
3.11e+05					
zipcode_98144	2.827e+05	7165.040	39.455	0.000	2.69e+05
2.97e+05					
zipcode_98146	1.294e+05	7342.058	17.621	0.000	1.15e+05
1.44e+05					
zipcode_98148	6.152e+04	1.25e+04	4.912	0.000	3.7e+04
8.61e+04					
zipcode_98155	1.554e+05	6586.017	23.588	0.000	1.42e+05
1.68e+05					
zipcode_98166	1.174e+05	7950.541	14.771	0.000	1.02e+05
1.33e+05					
zipcode_98168	6.001e+04	7585.424	7.912	0.000	4.51e+04
7.49e+04					
zipcode_98177	2.244e+05	7994.410	28.071	0.000	2.09e+05
2.4e+05					
zipcode_98178	7.431e+04	7442.250	9.985	0.000	5.97e+04
8.89e+04					

zipcode_98188	4.692e+04	9296.686	5.047	0.000	2.87e+04
6.51e+04					
zipcode_98198	4.784e+04	7434.966	6.435	0.000	3.33e+04
6.24e+04					
zipcode_98199	3.922e+05	7476.812	52.459	0.000	3.78e+05
4.07e+05					

```

=====
Omnibus:                1844.788    Durbin-Watson:                2.004
Prob(Omnibus):           0.000    Jarque-Bera (JB):            6380.863
Skew:                    0.559    Prob(JB):                     0.00
Kurtosis:                5.848    Cond. No.                     5.37e+05
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.37e+05. This might indicate that there are strong multicollinearity or other numerical problems.

"""

8 CONCLUSIONS & RECOMMENDATIONS

8.0.1 Key Takeaways

Our final model has an R2 value of 0.806, indicating that with the included variables, the model is capable of explaining 80.6% of the variability in a property's price.

As we can see in our three plots below, there does seem to be a strong linear relationship between price and our three selected variables: living space above ground, living space below ground and grade.

According to our model, for each foot of living space above ground that is increased, we see an increase in property value of approximately \$130.33. For each foot of living space below ground that is increased, we see an increase in property value of approximately \$90.98. Lastly, when the property grade is increased by 1 point, we see an increase in property value of approximately \$34,760.

An idea for future analysis would be to explore what costs would be involved in making these renovations, and to determine whether these recommendations would be cost-effective.

```

[50]: # Define function to notate dollar amounts in thousands
def thousands(x, pos):
    """Source: https://stackoverflow.com/questions/61330427/
    ↪set-y-axis-in-millions"""
    'The two args are the value and tick position'
    return '%1.0fK' % (x * 1e-3)

formatter = FuncFormatter(thousands)

```

```

[51]: # Group data by sqft_above and find aggregate mean
df_sqftabove = df_unscaled.groupby('sqft_above').mean()


[52]: # Reset index before plotting
df_sqftabove.reset_index(inplace=True)

[53]: # Plot scatter graph for sqft_above vs price
fig, ax = plt.subplots(figsize=(7,7))

sns.regplot(data=df_sqftabove, x='sqft_above', y='price', marker='x', \
            line_kws={"color": "deepskyblue"}, \
            scatter_kws={"color": "dodgerblue"}, ax=ax)

ax.set_title('Living Space Above Ground vs. Price')
ax.set_xlabel('Square Footage Above Ground (ft)')
ax.set_ylabel('Price ($)')
ax.yaxis.set_major_formatter(formatter);
ax.set_ylim([150000, 800000]);

```



```
[54]: # Group data points by sqft_basement and calculate aggregate mean
df_sqftbasement = df_unscaled.groupby('sqft_basement').mean()

df_sqftbasement.reset_index(inplace=True)

[55]: # Plot scatter graph for sqft_below vs price
fig, ax = plt.subplots(figsize=(7,7))

sns.regplot(data=df_sqftbasement, x='sqft_basement', y='price', marker='x', \
            line_kws={"color": "deepskyblue"}, \
            scatter_kws={"color": "dodgerblue"}, ax=ax)

ax.set_title('Living Space Below Ground vs. Price')
ax.set_xlabel('Square Footage Below Ground (ft)')
```

```
ax.set_ylabel('Price ($)')
ax.yaxis.set_major_formatter(formatter);
ax.set_ylim([300000, 700000]);
```



```
[56]: # Plot bar graph for grade vs price
fig, ax = plt.subplots(figsize=(7,7))

sns.barplot(data=df_unscaled, x='grade', y='price', palette='cool', ax=ax)

ax.set_title('House Grade vs. Price')
ax.set_xlabel('Grade (out of 13)')
ax.set_ylabel('Price ($)')
ax.yaxis.set_major_formatter(formatter);
```



[]: